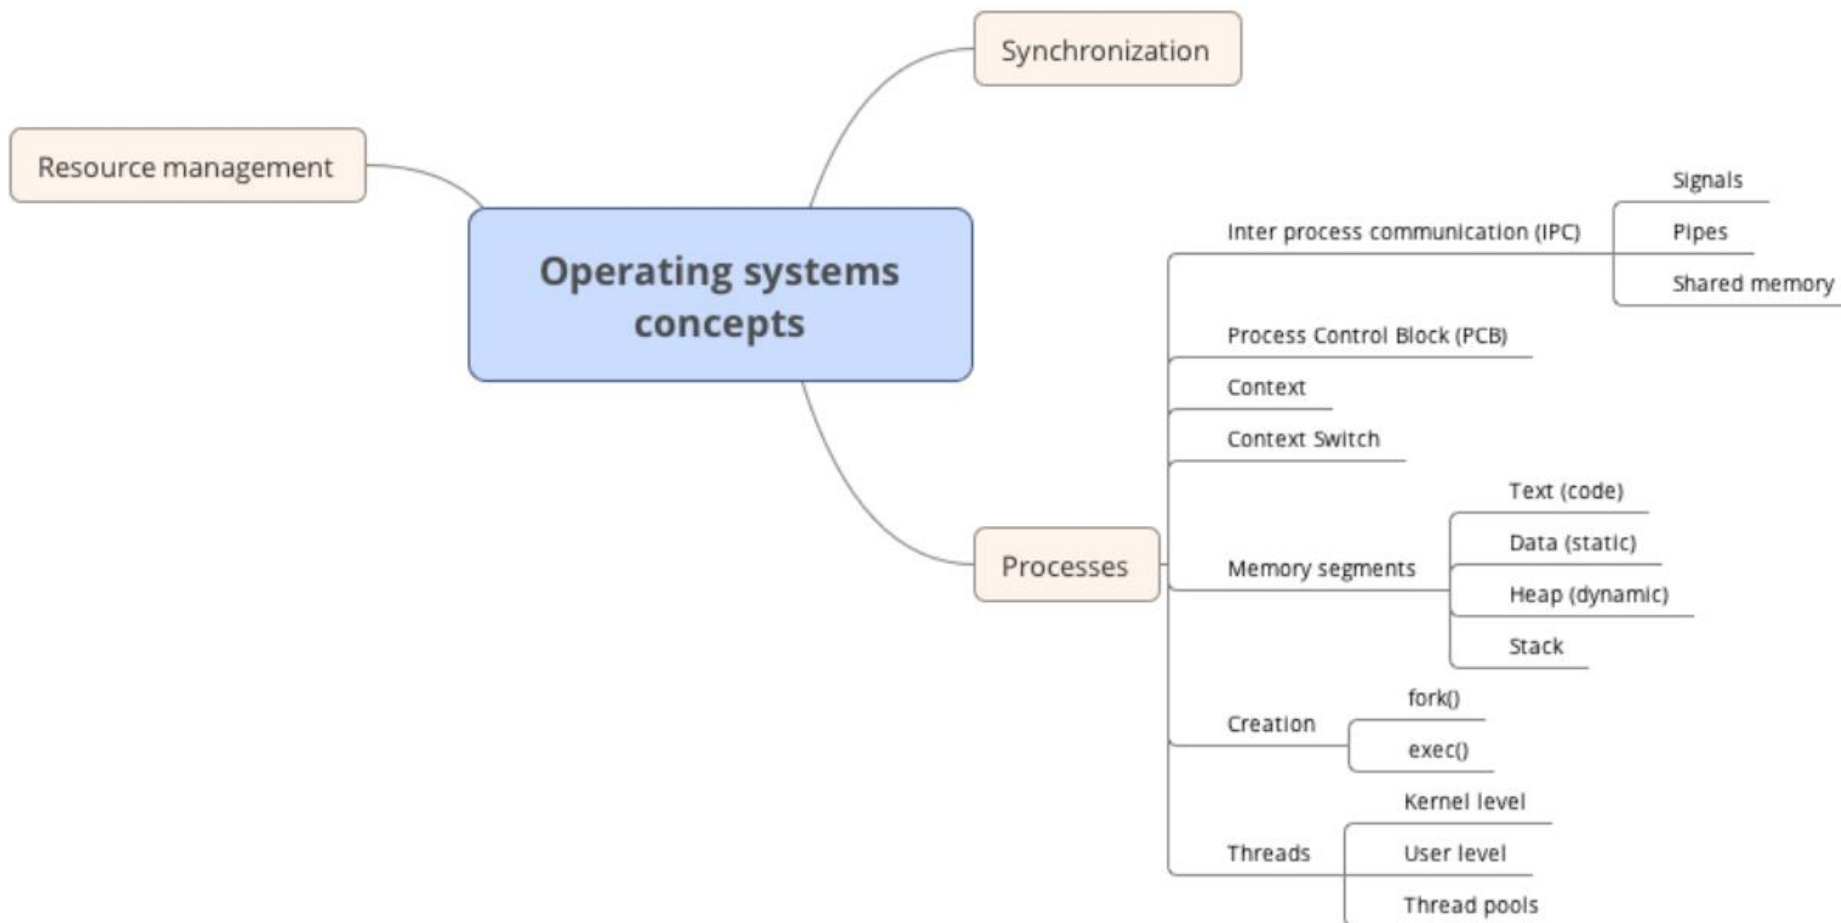



CS341: Operating System

Course Files:

<https://u.pcloud.link/publink/show?code=kZITAQXZCxtf3rGMIhpo90X7HMyGRXe3cBOK>

Process Management






Process control block

Process id (PID)
Process state (new, ready, running, waiting or terminated)
CPU Context
I/O status information
CPU scheduling information
Memory management information

Process memory image



From the Architecture to the OS to the User: Architectural, resources, OS management, and User Abstractions.

Hardware abstraction	Example OS Services	User abstraction
Processor	Process management, Scheduling, Traps, protection, accounting, synchronization	Process
Memory	Management, Protection, virtual memory	Address spaces
I/O devices	Concurrency with CPU, Interrupt handling	Terminal, mouse, printer, system calls
File System	File management, Persistence	Files
Distributed systems	Networking, security, distributed file system	Remote procedure calls, network file system

CPU starts...

CPU starts and loads instructions starting at 0xffffffff0



- Instruction jumps to BIOS code
- BIOS (Basic Input/Output System) is started
 - Performs basic tests (memory, keyboard, etc) – POST (power on self test)
 - Determines the “boot device” (Hard disk, Floppy, CD-ROM)
 - Loads the contents of the first physical sector (the Master Boot Record -

MBR - Cyl 0, Head 0, Sect 1) in memory 0x7C00 - 0x7DFF

- Jumps to 0x7C00
 - MBR code finds an “active” file system, loads the corresponding boot sector in memory, and jumps to it
 - The boot sector code loads the *operating system*

Concurrency and Parallelism

- Concurrent multithreading systems **give the appearance of several tasks executing at once**, but these tasks are actually split up into chunks that share the processor with chunks from other tasks.
- In parallel systems, two tasks are actually performed simultaneously. **Parallelism requires a multi-CPU system.**

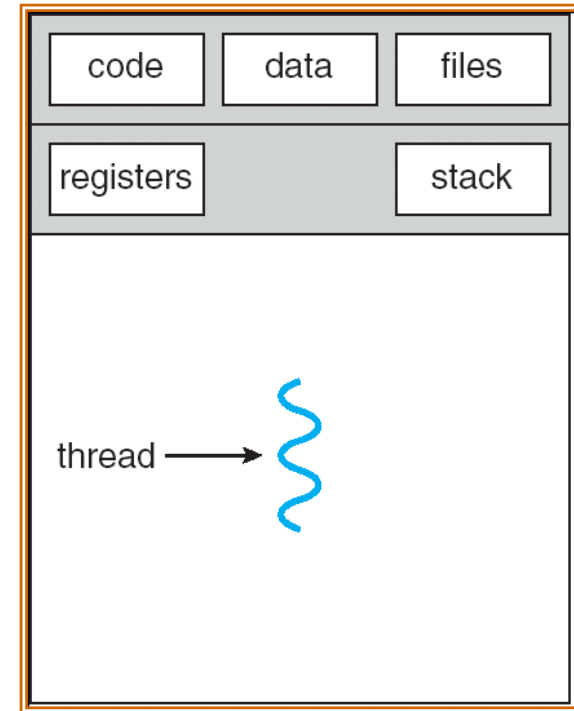


Process Management

- A process as the unit of execution.
- How are processes represented in the OS?
- What are possible execution states and how does the system move from one state to another?
- How are processes created in the system?
- How do processes communicate? Is this efficient?

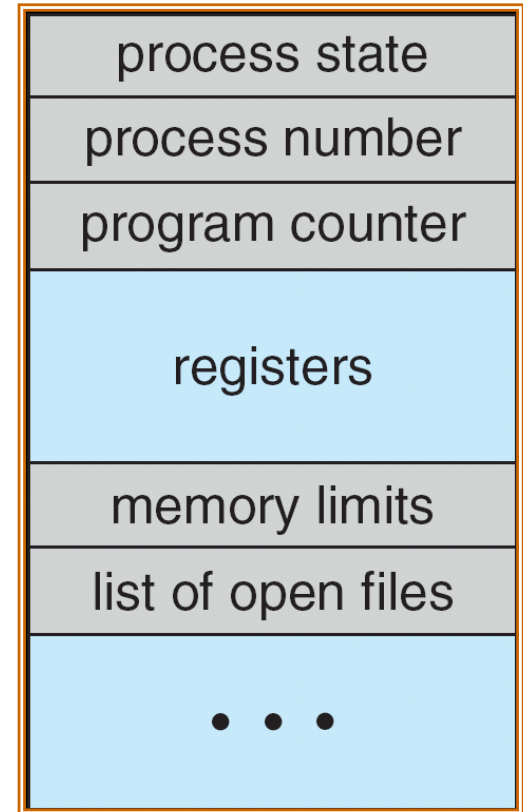
Starting Point: Single Threaded Process

- Process: OS abstraction of what is needed to run a single program
 1. Sequential program execution stream
 - Sequential stream of execution (thread)
 - State of CPU registers
 2. Protected resources
 - Contents of Address Space
 - I/O state



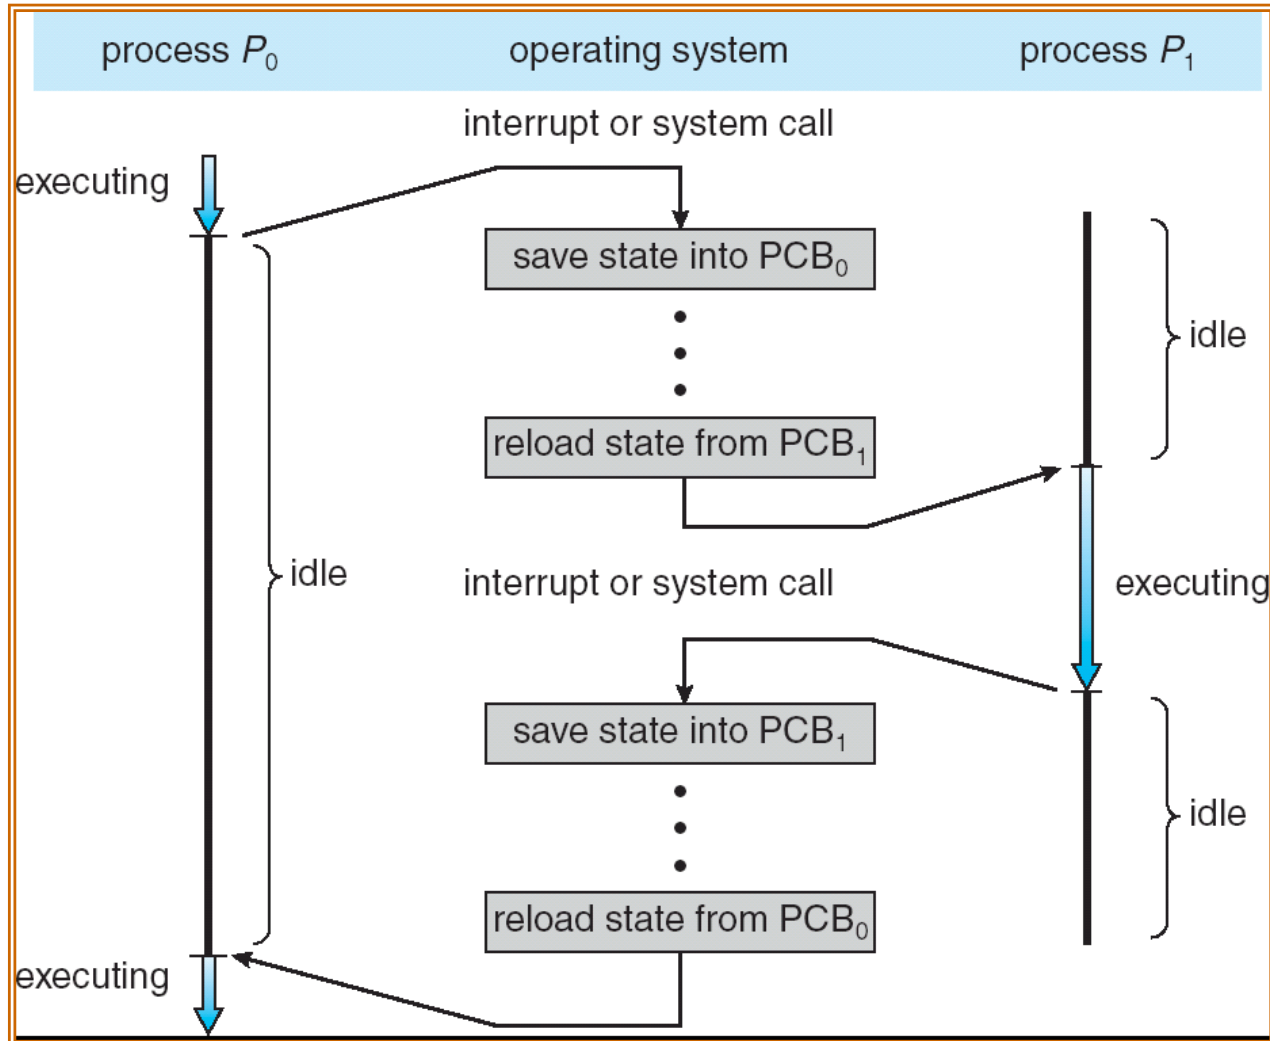
Multiplexing Processes

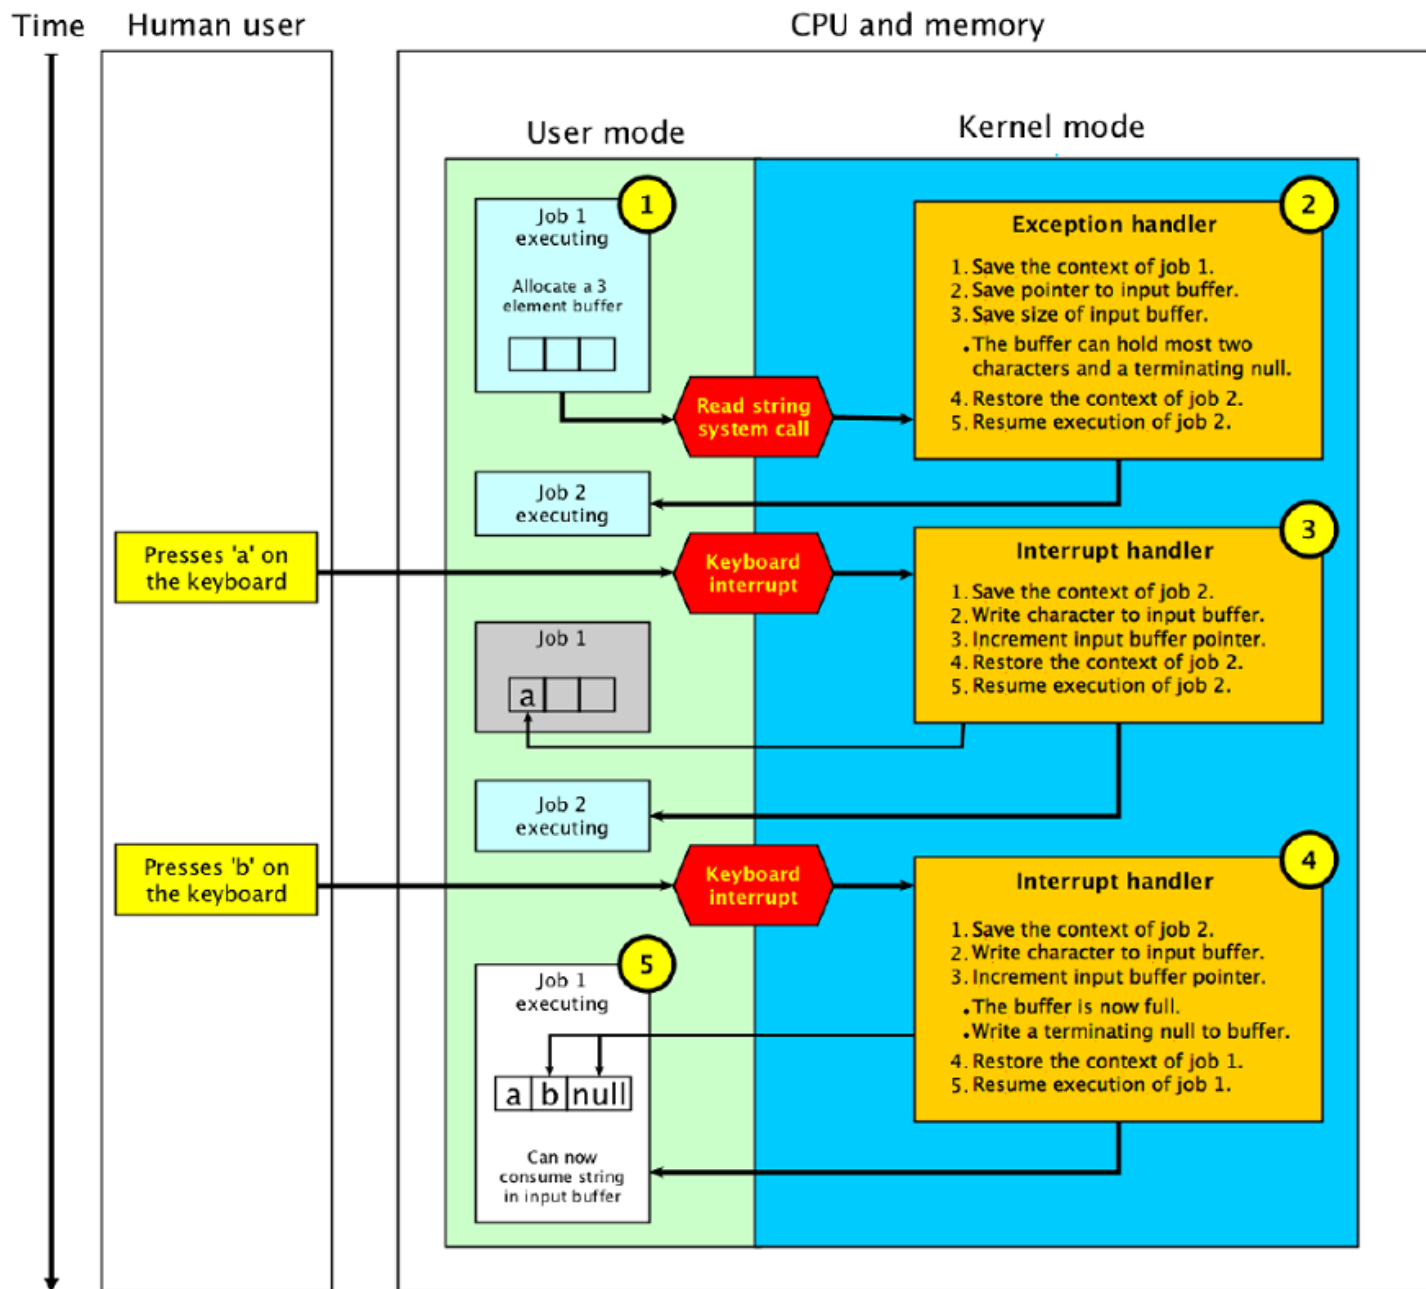
- Snapshot of each process in its PCB
 - Only one active at a time (per core...)
- Give out CPU to different processes
 - **Scheduling**
 - **Policy Decision**
- Give out non-CPU resources
 - Memory/IO
 - Another **policy decision**



Process
Control
Block

Context Switch





Context Switch

Starting and stopping processes is called a **context switch**, and is a relatively expensive operation.

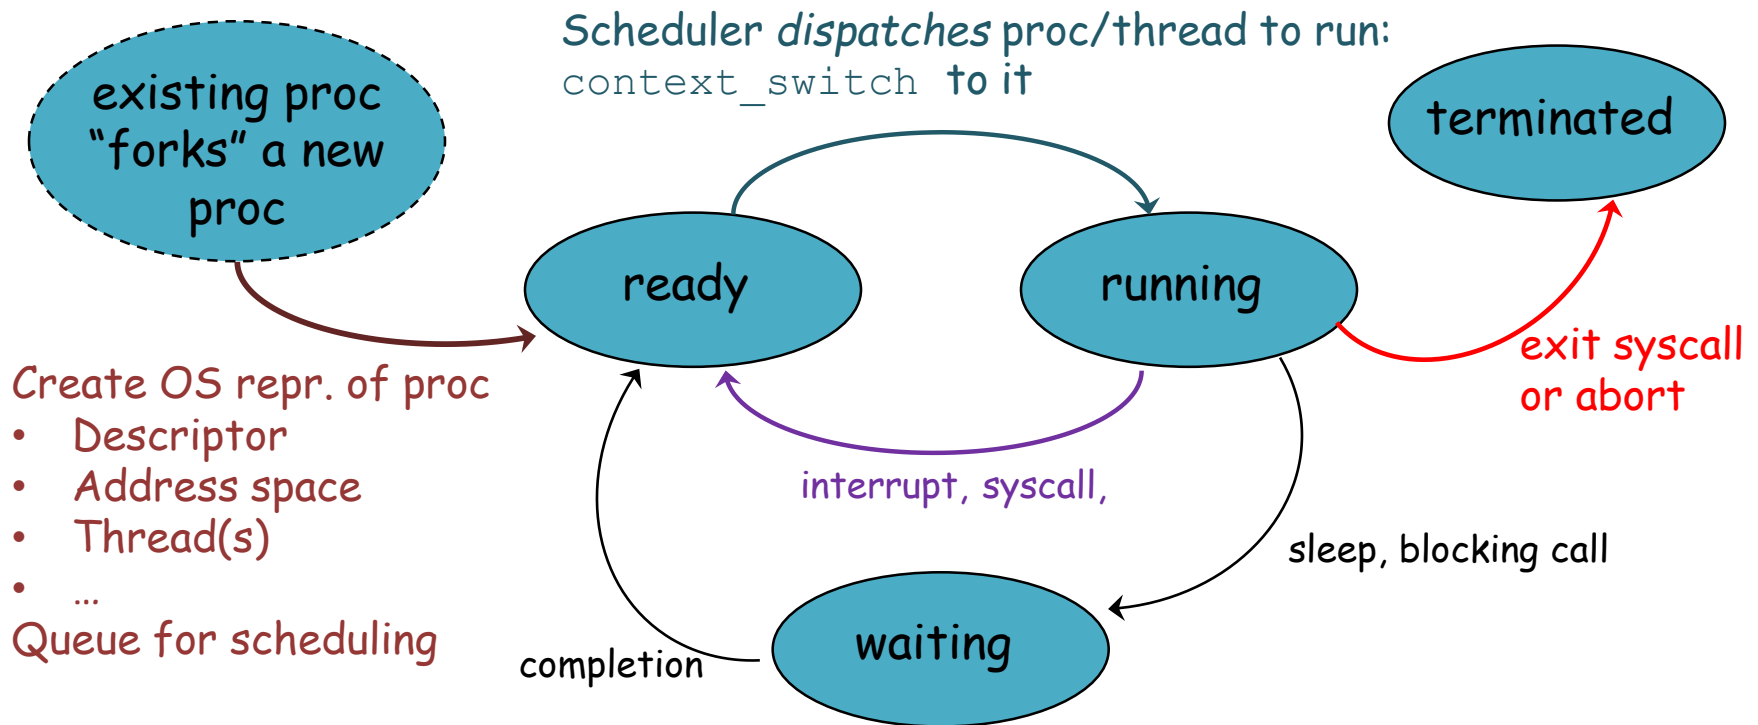
- The OS starts executing a ready process by loading hardware registers (PC, SP, etc) from its PCB
- While a process is running, the CPU modifies the Program Counter (PC), Stack Pointer (SP), registers, etc.
- When the OS stops a process, it saves the current values of the registers, (PC, SP, etc.) into its PCB
- This process of switching the CPU from one process to another (stopping one and starting the next) is the context switch.
 - Time sharing systems may do 100 to 1000 context switches a second.
 - The cost of a context switch and the time between switches are closely related

Creating a Process

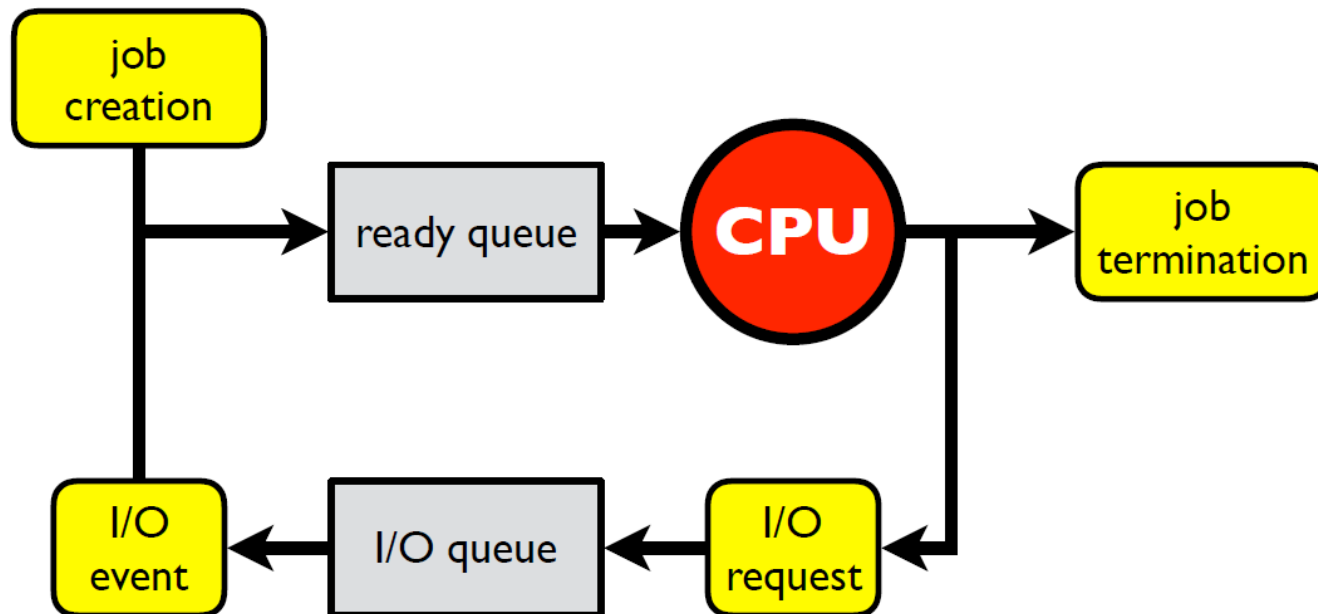
One process can create other processes to do work.

- The creator is called the *parent* and the new process is the *child*
 - The parent defines (or donates) resources and privileges to its children
 - A parent can either wait for the child to complete, or continue in parallel
 - In Unix, the *fork* system call is used to create child processes
 - Fork copies variables and registers from the parent to the child
 - The *only difference* between the child and the parent is the value returned by fork
-
- * In the parent process, fork returns the process id of the child
 - * In the child process, the return value is 0
 - The parent can wait for the child to terminate by executing the *wait* system call or continue execution
 - The child often starts a new and different program within itself, via a call to *exec* system call.

Lifecycle of a process



- OS juggles many process/threads using kernel data structures
- Proc's may create other process (fork/exec)
 - All starts with init process at boot



Process Execution State

Execution state of a process indicates what it is doing

new: the OS is setting up the process state

running: executing instructions on the CPU


ready: ready to run, but waiting for the CPU

waiting: waiting for an event to complete

terminated: the OS is destroying this process

- As the program executes, it moves from state to state, as a result of the program actions (e.g., system calls), OS actions (scheduling), and external actions (interrupts).

Process Execution State



state sequence

new

ready

running

waiting for I/O

ready

running

terminated

– Example:



```
– void main() {  
  printf('Hello World');  
}
```

- The OS manages multiple active process using *state queues* (More on this in a minute...)

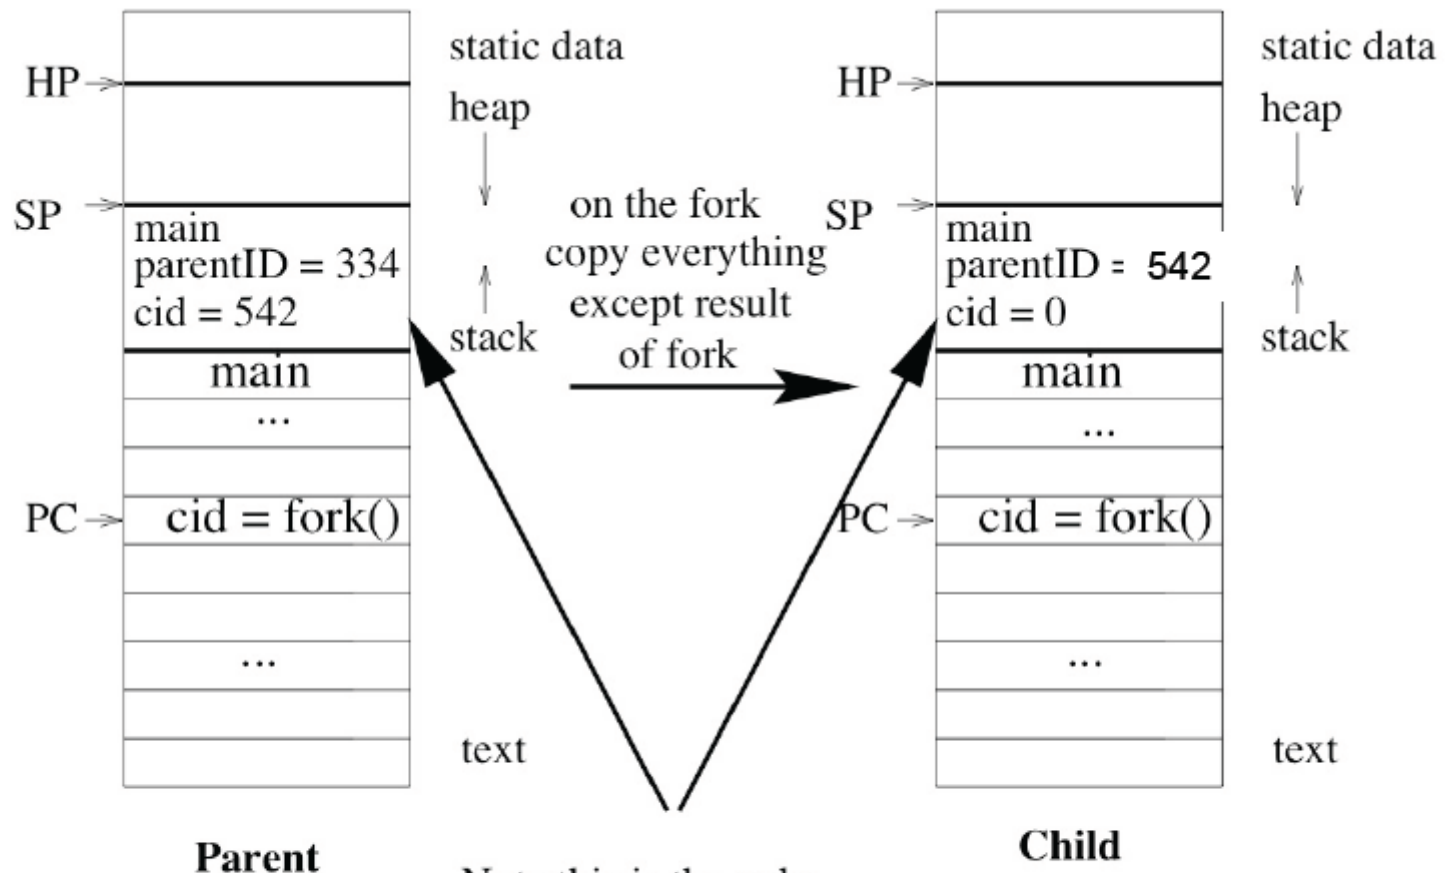
Creating a Process: Example

When you log in to a machine running Unix, you create a shell process.

- Every command you type into the shell is a child of your shell process and is an implicit *fork* and *exec* pair.
- For example, you type *emacs*, the OS "*forks*" a new process and then "*exec*" (executes) *emacs*.
- If you type an *&* after the command, Unix will run the process in parallel with your shell, otherwise, your next shell command must wait until the first one completes.



What is happening on the Fork



Note this is the only difference between the parent and the child at the time of the fork.

Example Unix Program: Explanation


fork() forks a new child process that is a copy of the parent.

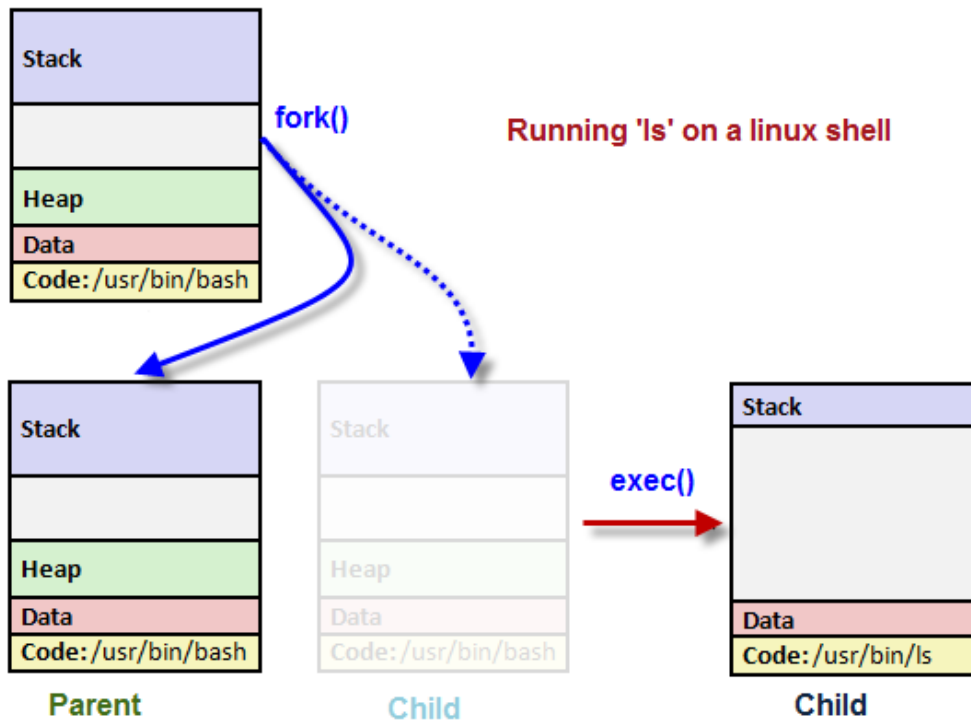
execvp() replaces the program of the current process with the named program.

sleep() suspends execution for at least the specified time.

waitpid() waits for the named process to finish execution.

gets() reads a line from a file.





Example Program: Fork

```
#include <stdio.h>
#include <unistd.h>

int main(int argc, char ** argv)
{
    int pid;

    printf("Before forking in the parent\n");

    pid = fork();
    if (pid == 0)
        printf("I am the child.\n");
    else
        printf("I am the proud parent of pid %d\n", pid);

    return 0;
}
```

Example Program: Fork

```
#!/usr/bin/env python2

import os

print 'Before forking in the parent'

pid = os.fork()
if pid == 0:
    print 'I am the child'
else:
    print 'I am the proud parent of pid', pid
```

Example Program: Fork

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char ** argv)
{
    int pid;

    pid = fork();
    if (pid == 0)
        execl("/bin/echo", "/bin/echo", "hello", (char *) NULL);

    printf("Only the parent gets here\n");

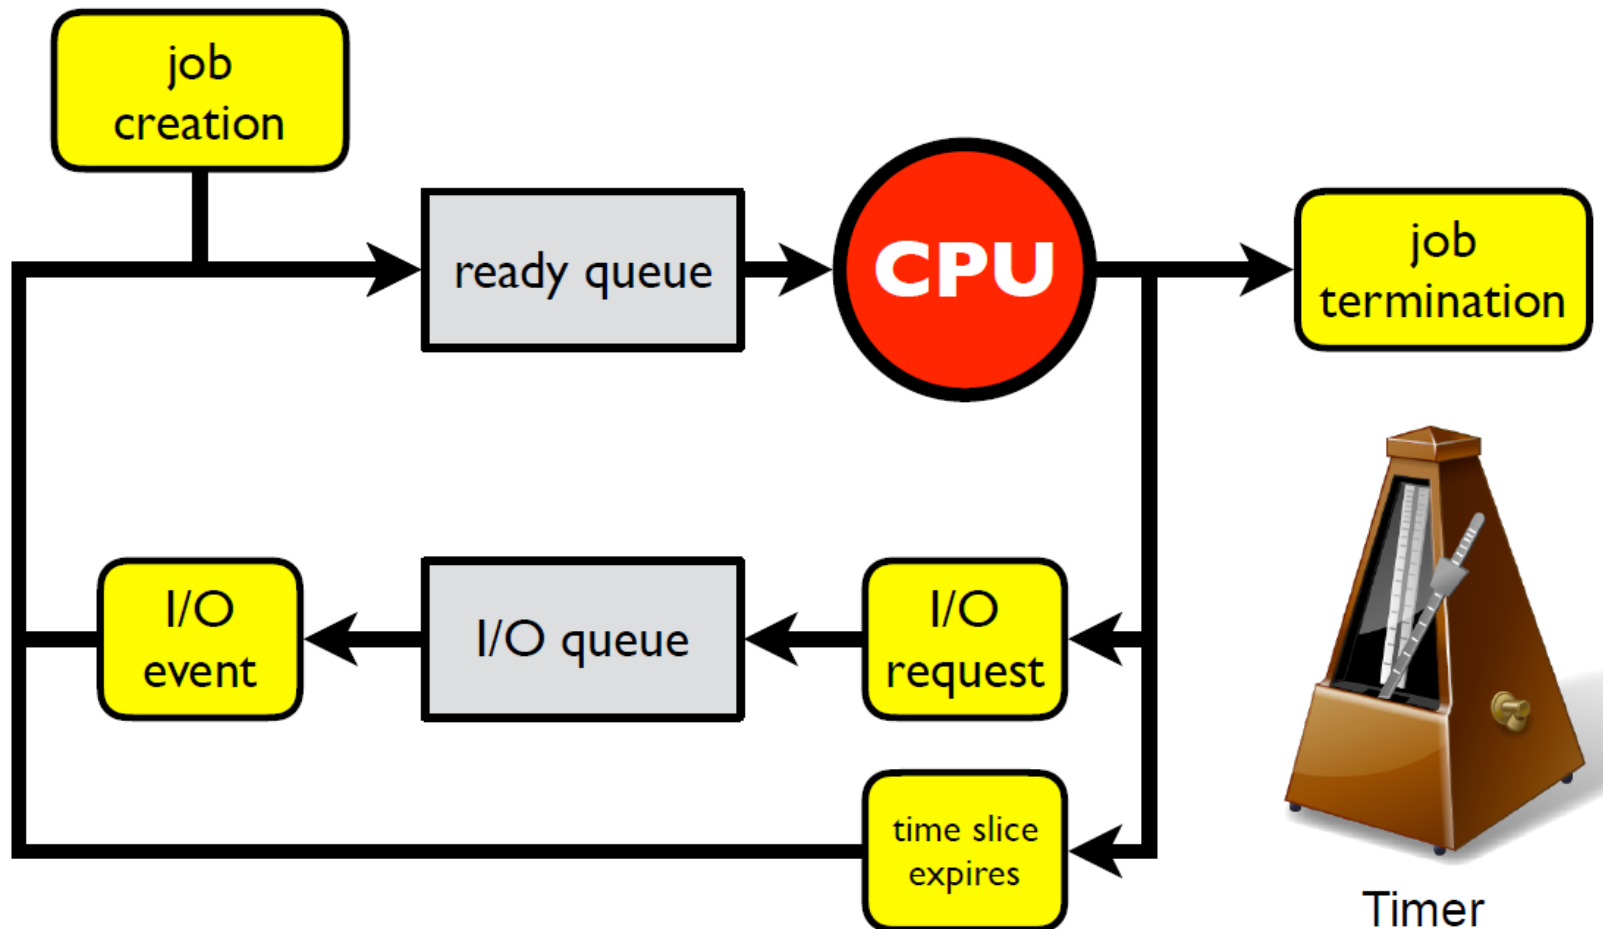
    return 0;
}
```

Example Unix Program: Fork

```
#include <sys/wait.h>
#include <stdio.h>
main() {
    int parentID = getpid(); /* ID of this process */
    char prgname[1024];
    gets(prgname); /* read the name of program we want to start */
    int cid = fork();
    if(cid == 0) { /* I'm the child process */
        execlp( prgname, prgname, 0); /* Load the program */
        /* If the program named prgname can be started, we never get
        to this line, because the child program is replaced by prgname */
        printf("I didn't find program %s\n", prgname);
    } else { /* I'm the parent process */
        sleep (1); /* Give my child time to start. */
        waitpid(cid, 0, 0); /* Wait for my child to terminate. */
        printf("Program %s finished\n", prgname);
    } }
```

Function	Specification of program file (\neg , p)	Specification of arguments (v , l)	Source of environment (e , $-$)
<i>execve()</i>	pathname	array	<i>envp</i> argument
<i>execle()</i>	pathname	list	<i>envp</i> argument
<i>execlp()</i>	filename + PATH	list	caller's <i>environ</i>
<i>execvp()</i>	filename + PATH	array	caller's <i>environ</i>
<i>execv()</i>	pathname	array	caller's <i>environ</i>
<i>execl()</i>	pathname	list	caller's <i>environ</i>

A schematic view of multitasking



A schematic view of multiprogramming

