


# Inter-Process Communication




How can we make  
processes communicate  
(sharing information)?

Two methods: **message passing** and **shared memory**.

## Message passing an shared memory



Communication take place by means of messages exchanged between the cooperating processes.



A region of memory that is shared by cooperating process is established. Process can then exchange information by reading and writing data to the shared region.

# Message passing

Messages can be exchanged between processes either directly or indirectly using a common mailbox.

The recipient process usually must give its permission for communication to take place with an accept connection system call.

Message-passing is **useful** for exchanging **smaller amounts of data**, because no conflicts need be avoided.

**Easier to implement** compared to the shared memory approach.

# shared memory


Processes can communicate by reading and writing to shared memory.

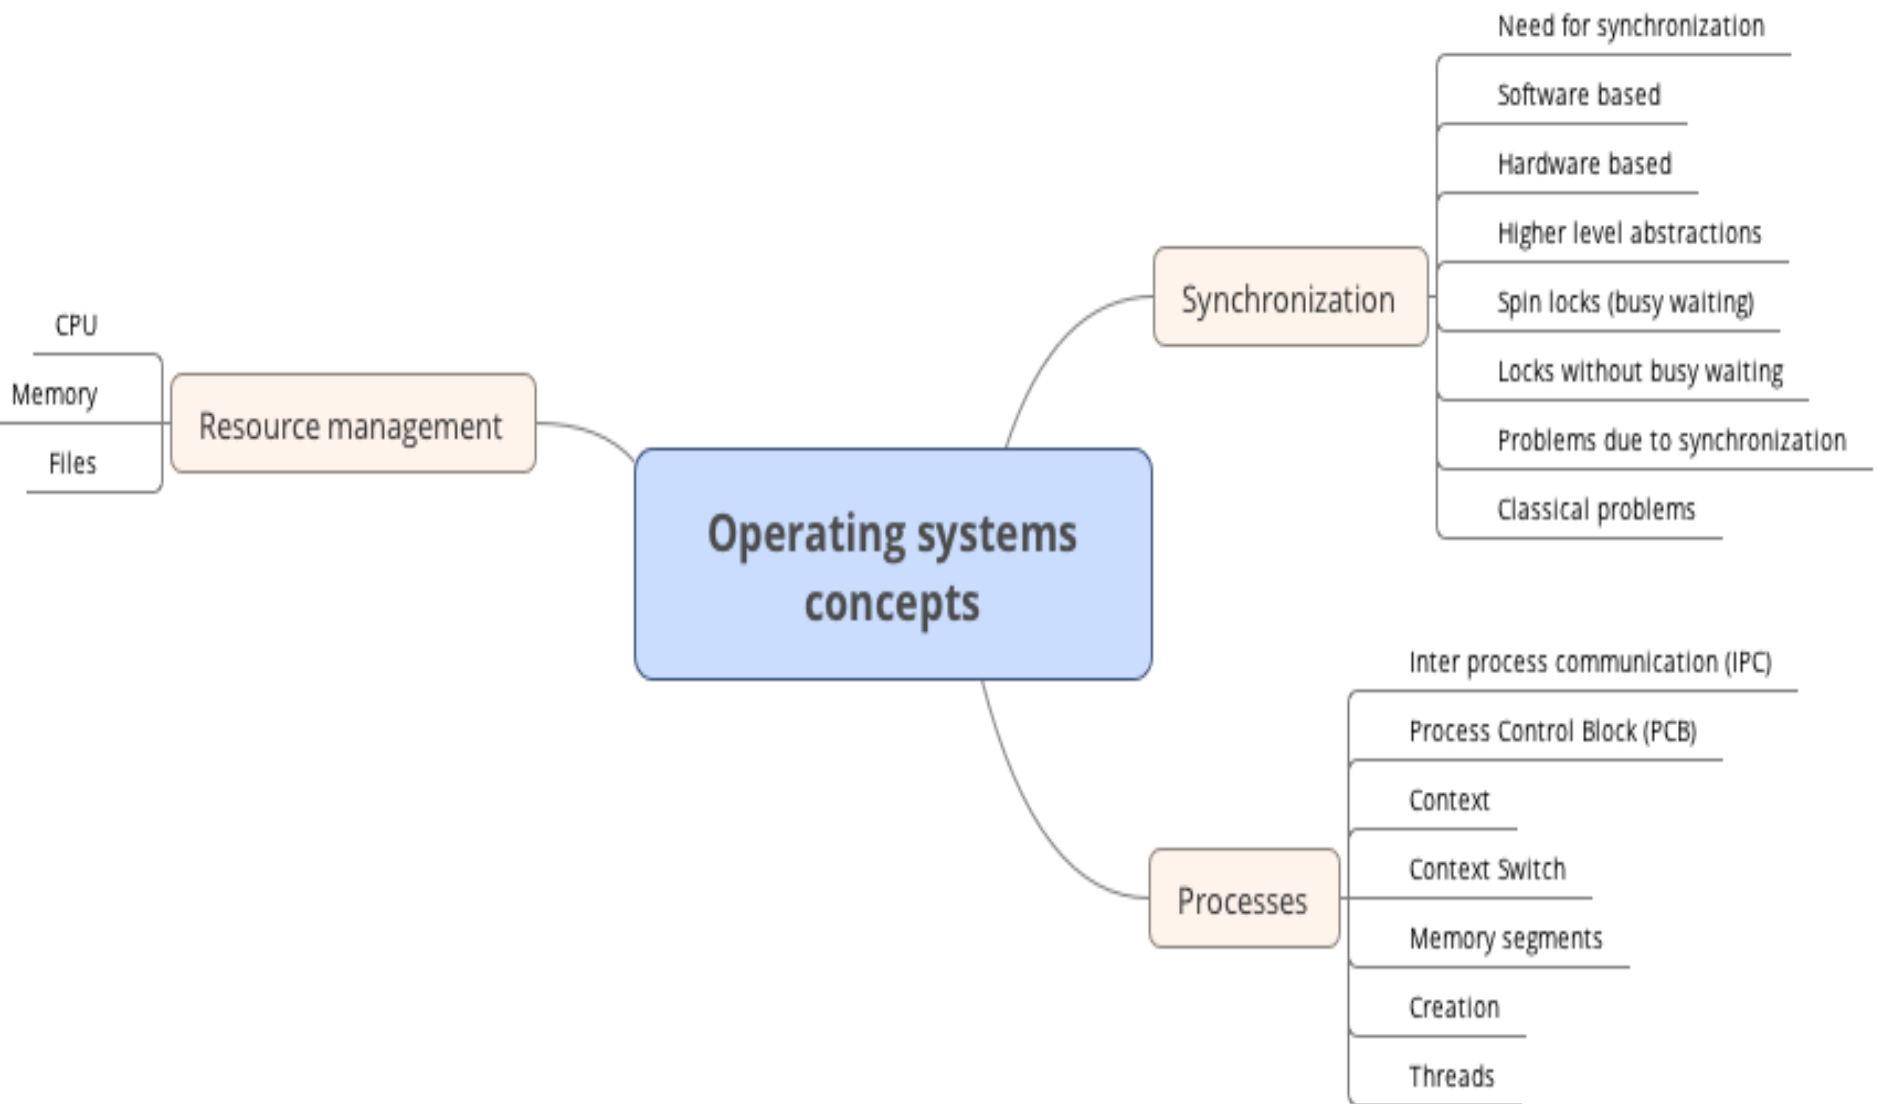
Normally, the OS tries to prevent one process from accessing another process's memory. Shared-memory requires that two or more processes agree to remove this restriction.

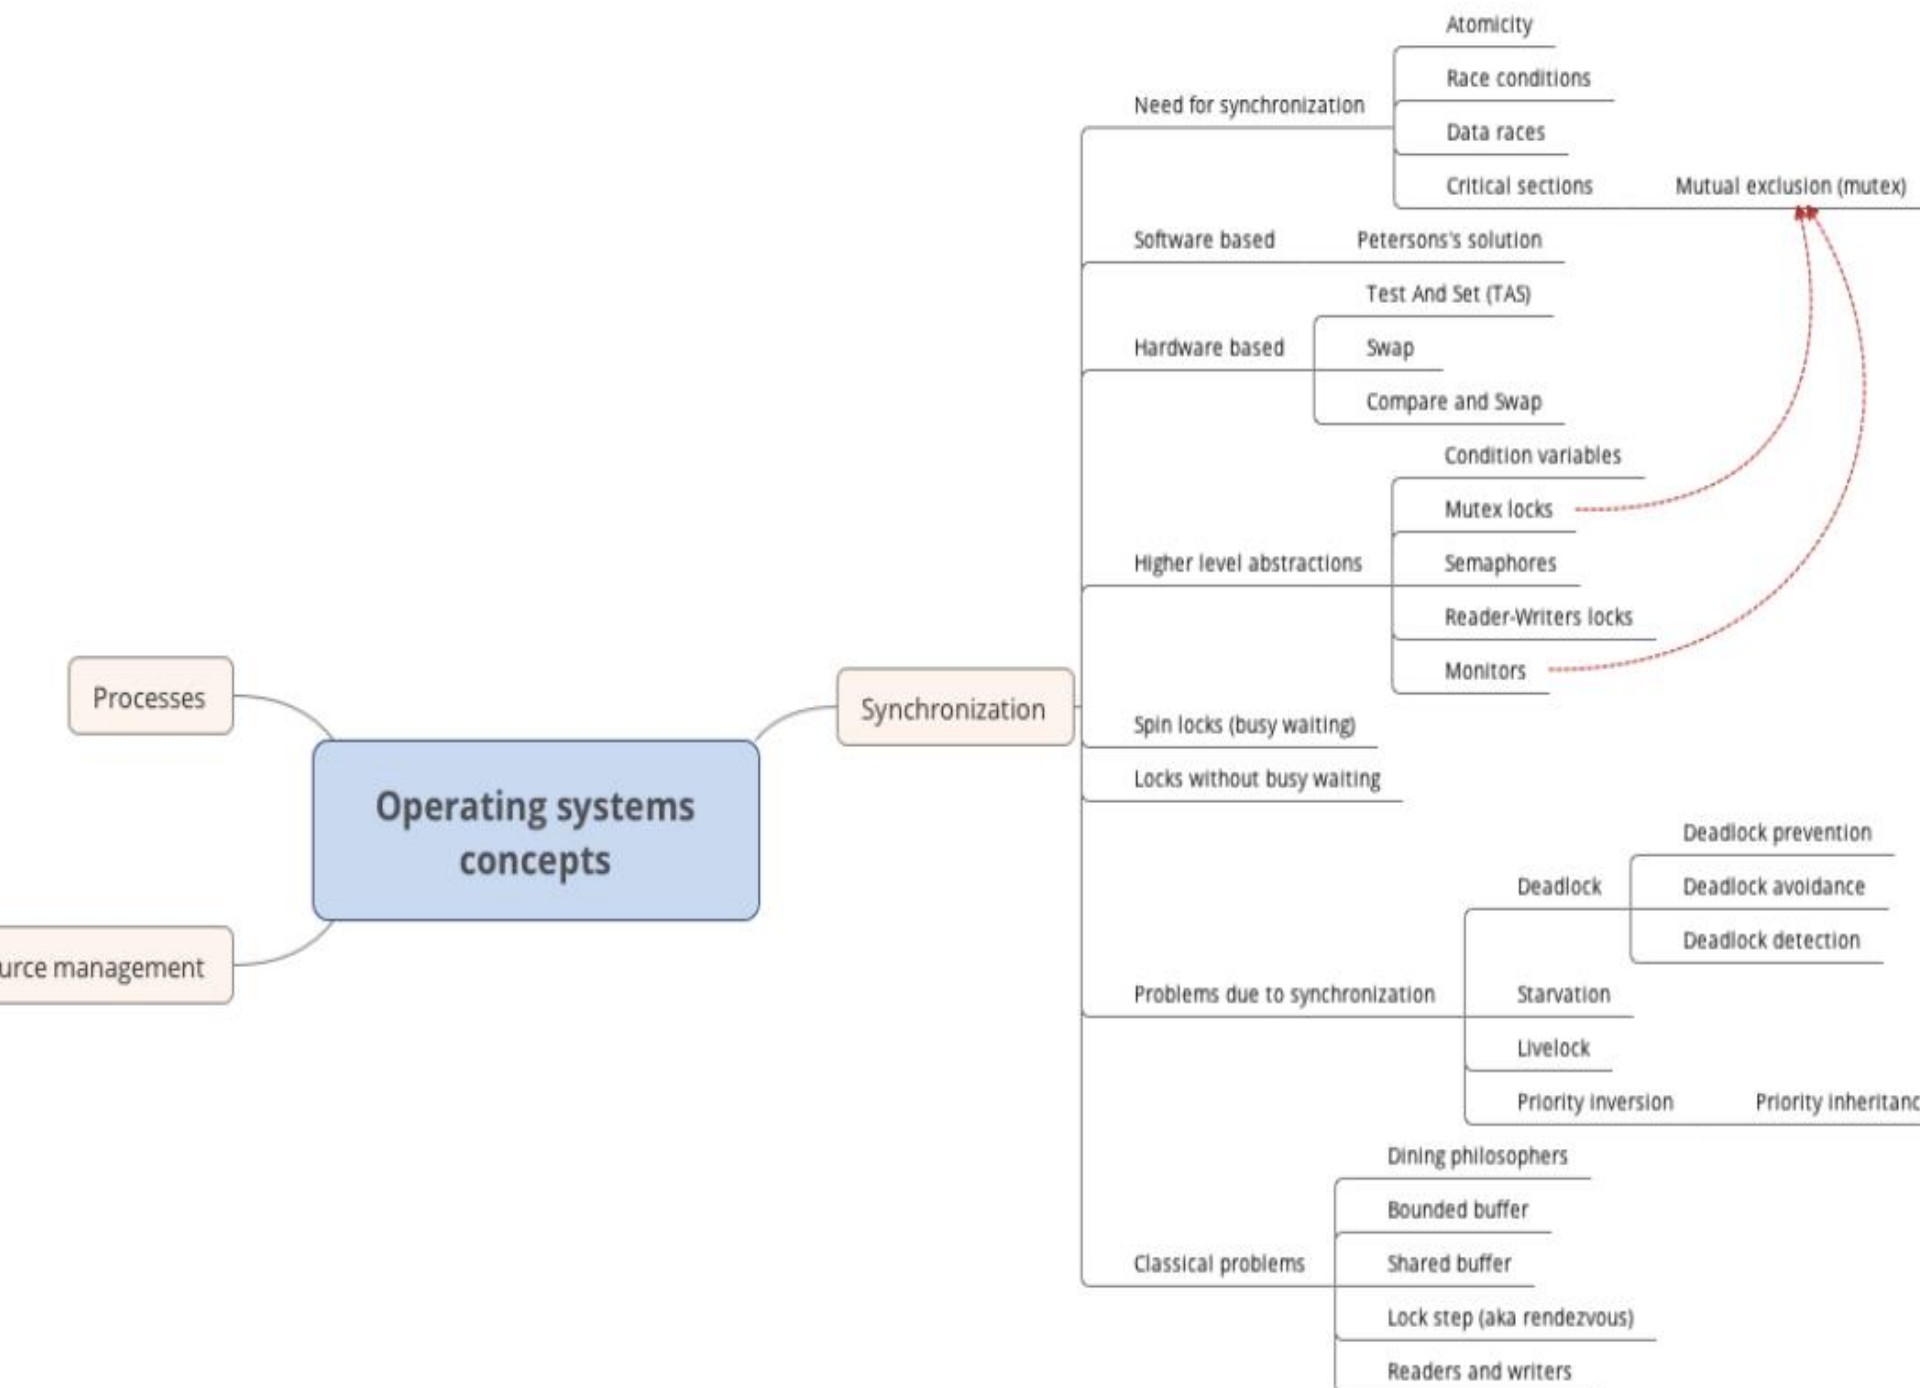
Shared memory **allows maximum speed** and convenience of communication, since it can be done at memory transfer speeds.

**Problems:** protection and synchronization between the processes sharing memory.


# Concurrency/ Synchronization/ Deadlocks







Large problems can often be divided into smaller ones, which are then solved in **parallel** (at the same time) on multiple physical CPUs.

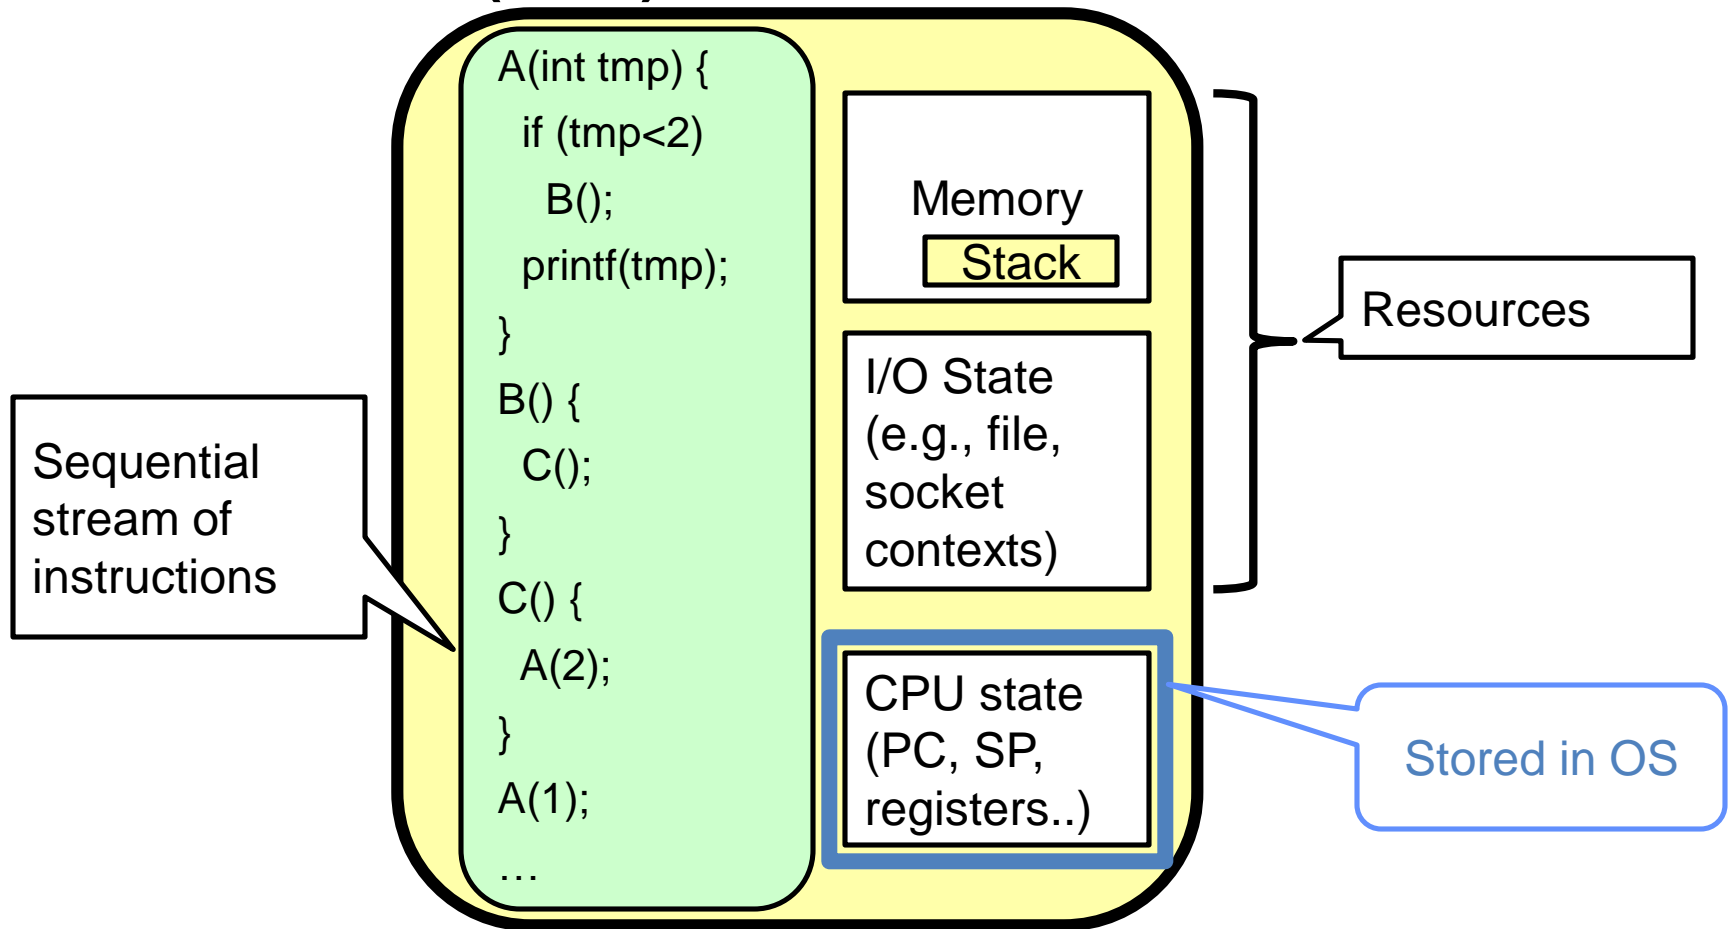


This is not the same as multithreading. Multithreading can be done on a single core CPU. In such a case, two threads can never execute at the same time on the CPU.

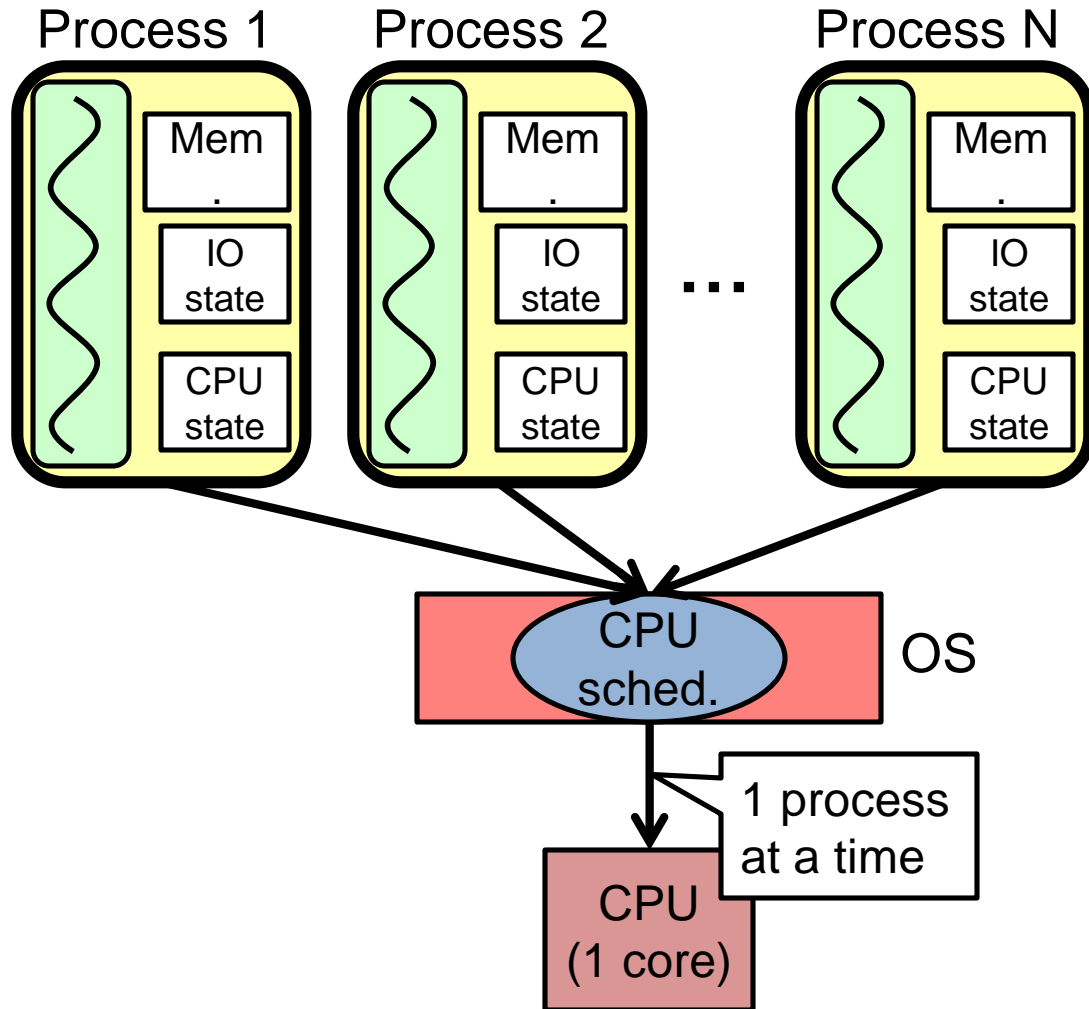


# Process

## (Unix) Process

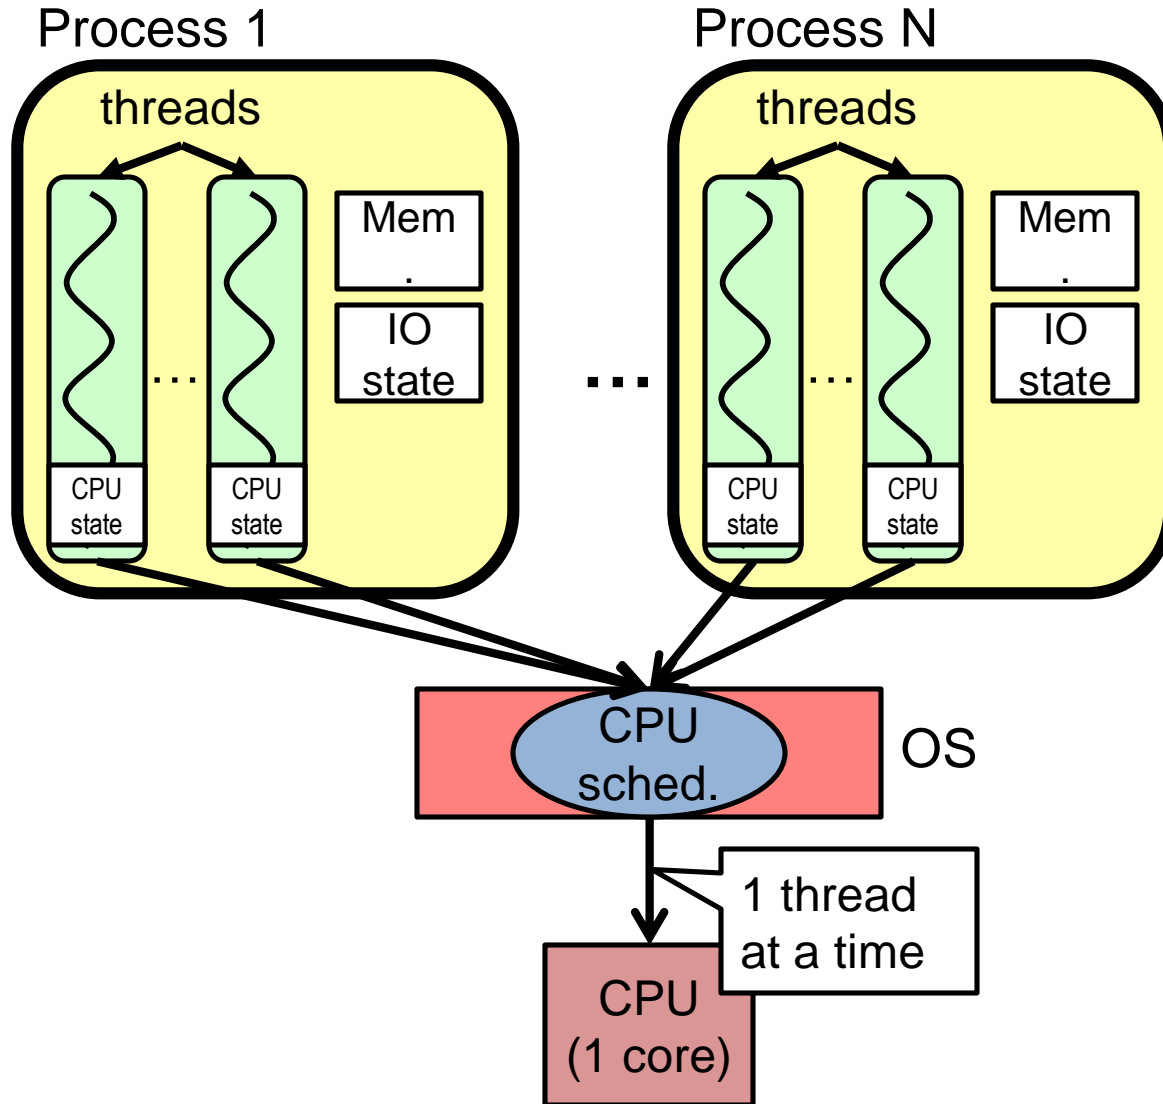


# Putting it together: Processes



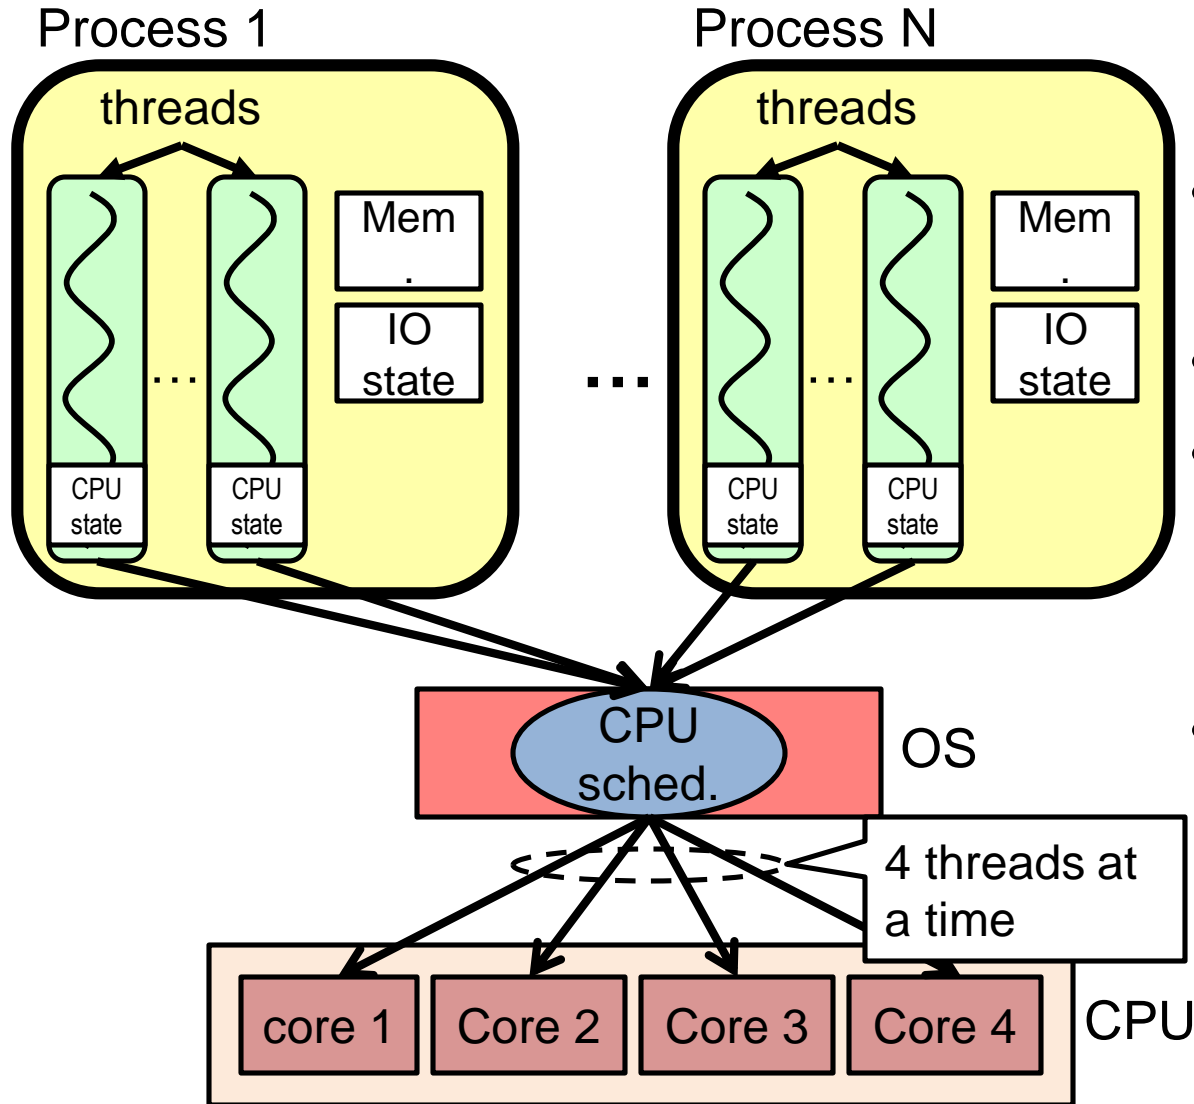
- Switch overhead: **high**
  - CPU state: **low**
  - Memory/IO state: **high**
- Process creation: **high**
- Protection
  - CPU: **yes**
  - Memory/IO: **yes**
- Sharing overhead: **high** (involves at least a context switch)

# Putting it together: Threads



- Switch overhead: **low** (only CPU state)
- Thread creation: **low**
- Protection
  - CPU: **yes**
  - Memory/IO: **No**
- Sharing overhead: **low** (thread switch overhead low)


# Putting it together: Multi-Cores



- Switch overhead: **low** (only CPU state)
- Thread creation: **low**
- Protection
  - CPU: **yes**
  - Memory/IO: **No**
- Sharing overhead: **low** (thread switch overhead low)

# Memory Footprint of Two-Thread Example

- If we stopped this program and examined it with a debugger, we would see
  - Two sets of CPU registers
  - Two sets of Stacks
- Questions:
  - How do we position stacks relative to each other?
  - What maximum size should we choose for the stacks?
  - What happens if threads violate this?
  - How might you catch violations?



# Concurrency

The ability of different parts or units of a program, algorithm, or problem to be executed out-of-order or in partial order, without affecting the final outcome.

This allows for parallel execution of the concurrent units, which can significantly improve overall speed of the execution in multi-processor and multi-core systems.

# Concurrency

≠

# Parallelism


Concurrency is often referred to as the decomposability property of a program, algorithm, or problem into order-independent or partially-ordered components or units.

**Source:** Lamport, Leslie (July 1978). "Time, Clocks, and the Ordering of Events in a Distributed System"

# Concurrent execution of threads





## On a single core CPU



Threads take turn executing on the single CPU core. By switching fast enough between the threads they appear to be executing "at the same time".

## On a dual core CPU




concurrent  $\neq$  parallel







# Concurrency and response time



Creating a new process is time consuming and resource intensive.

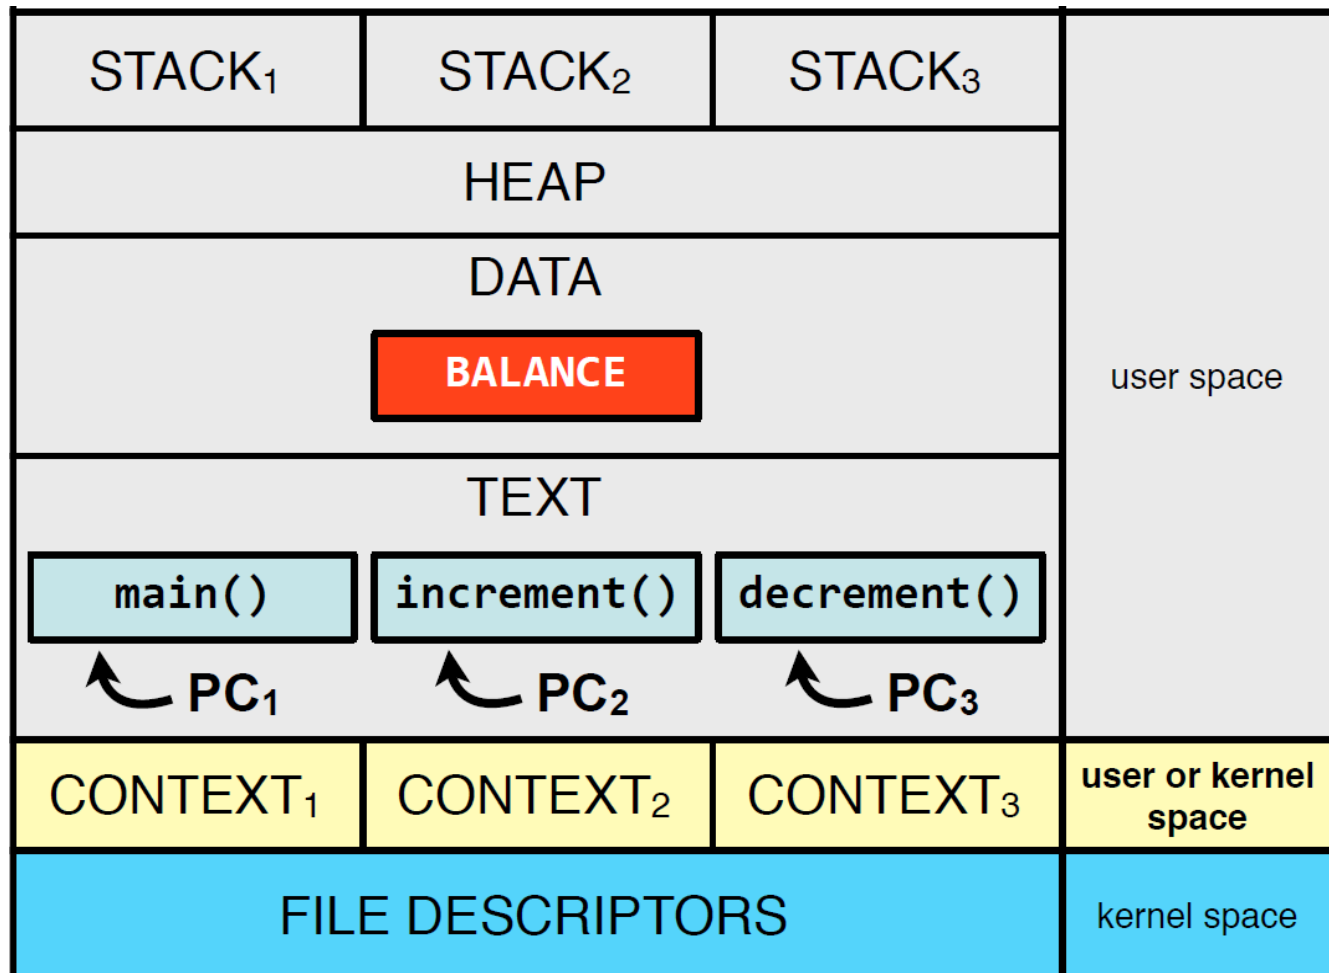


# ATM Bank Server



- ATM server problem:
  - Service a set of requests
  - Do so without corrupting database
  - Don't hand out too much money

# A process with three threads



**main()**

```
#define N 10000  
int BALANCE 0;
```

Thread A

**increment()**

```
for (int i = 0; i < N; i++) {  
    BALANCE++;  
}
```

Thread B

**decrement()**

```
for (int i = 0; i < N; i++) {  
    BALANCE--;  
}
```

**BALANCE == ?**

# Correctness with Concurrent Threads

- Non-determinism:
  - Scheduler can run threads in **any order**
  - Scheduler can switch threads **at any time**
  - This can make testing very difficult
- *Independent Threads*
  - No state shared with other threads
  - Deterministic, reproducible conditions
- *Cooperating Threads*
  - Shared state between multiple threads
- **Goal: Correctness by Design**

## Problem is at the lowest level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

$x = 1;$

Thread B

$y = 2;$

- However, What about (Initially,  $y = 12$ ):

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

- What are the possible values of  $x$ ?

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2$

$x = 13$

# Problem is at the lowest level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

$x = 1;$

Thread B

$y = 2;$

- However, What about (Initially,  $y = 12$ ):

Thread A

$x = 1;$

$x = y + 1;$

Thread B

$y = 2;$

$y = y * 2;$

- What are the possible values of  $x$ ?

Thread A

Thread B

$y = 2;$

$y = y * 2;$

$x = 1;$

$x = y + 1;$

$x = 5$



# Problem is at the lowest level

Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A

`x = 1;`

Thread B

`y = 2;`

However, What about (Initially,  $y = 12$ ):

Thread A

`x = 1;`

`x = y+1;`

Thread B

`y = 2;`

`y = y*2;`

– What are the possible values of  $x$ ?

Thread A

`x = 1;`

`x = y+1;`

Thread B

`y = 2;`

`y = y*2;`

`x=3`

# Correctness Requirements

- Threaded programs must work for all interleavings of thread instruction sequences
  - Cooperating threads inherently non-deterministic and non-reproducible
  - Really hard to debug unless carefully designed!
- Example: Therac-25
  - Machine for radiation therapy
    - Software control of electron accelerator and electron beam/Xray production
    - Software control of dosage
  - Software errors caused overdoses and the death of several patients
    - A series of race conditions on shared variables and poor software design
    - “They determined that data entry speed during editing was the key factor in producing the error condition: If the prescription data was edited at a fast pace, the overdose occurred.”





Figure 1. Typical Therac-25 facility

# Space Shuttle Example

- Original Space Shuttle launch aborted 20 minutes before scheduled launch
- Shuttle has five computers:
  - Four run the “Primary Avionics Software System” (PASS)
    - Asynchronous and real-time
    - Runs all of the control systems
    - Results synchronized and compared 440 times per second
  - The Fifth computer is the “Backup Flight System” (BFS)
    - Stays synchronized in case it is needed
    - Written by completely different team than PASS
- Countdown aborted because BFS disagreed with PASS
  - A 1/67 chance that PASS was out of sync one cycle
  - Bug due to modifications in **initialization** code of PASS
    - A delayed init request placed into timer queue
    - As a result, timer queue not empty at expected time to force use of hardware clock
  - Bug not found during extensive simulation



# Atomic Operations

- To understand a concurrent program, we need to know what the underlying atomic operations are!
- **Atomic Operation**: an operation that always runs to completion or not at all
  - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
  - Fundamental building block – if no atomic operations, then have no way for threads to work together
- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic
- Many instructions are not atomic
  - Double-precision floating point store often not atomic
  - VAX and IBM 360 had an instruction to copy a whole array

# Atomic Operations

- `uint64_t sharedValue = 0;`
- `void storeValue()`
  - `{`
  - `sharedValue = 0x100000002;`
  - `}`
- `$ gcc -O2 -S -masm=intel test.c`
- `$ cat test.s ...`
  - `mov DWORD PTR sharedValue, 2`
  - `mov DWORD PTR sharedValue+4, 1`
  - `ret`

# Concurrency Challenges

- Multiple computations (threads) executing in parallel to
  - share resources, and/or
  - share data
- Fine grain sharing:
  - ↑ increase concurrency → better performance
  - ↓ more complex
- Coarse grain sharing:
  - ↑ Simpler to implement
  - ↓ Lower performance
- Examples:
  - Sharing CPU for 10ms vs. 1min
  - Sharing a database at the row vs. table granularity

# Concurrency poses challenges for:


- **Correctness**
  - Threads accessing shared memory should not interfere with each other
- **Liveness**
  - Threads should not get stuck, should make forward progress
- **Efficiency**
  - Program should make good use of available computing resources (e.g., processors).
- **Fairness**
  - Resources apportioned fairly between threads

# Definitions

- **Synchronization**: using atomic operations to ensure cooperation between threads
  - For now, only loads and stores are atomic
- **Critical Section**: piece of code that only one thread can execute at once
- **Mutual Exclusion**: ensuring that only one thread executes critical section
  - One thread *excludes* the other while doing its task
  - Critical section and mutual exclusion are two ways of describing the same thing



# Critical Section Problem



Mutual exclusion using critical regions

# Race Conditions

- What are the possible values of **x** below?
- Initially **x = y = 0**;

Thread A

**x = 1;**

Thread B

**y = 2;**

- Must be **1**. Thread B cannot interfere.

# Race Conditions

- What are the possible values of  $x$  below?
- Initially  $x = y = 0$ ;

Thread A

$x = y + 1$ ;

Thread B

$y = 2$ ;

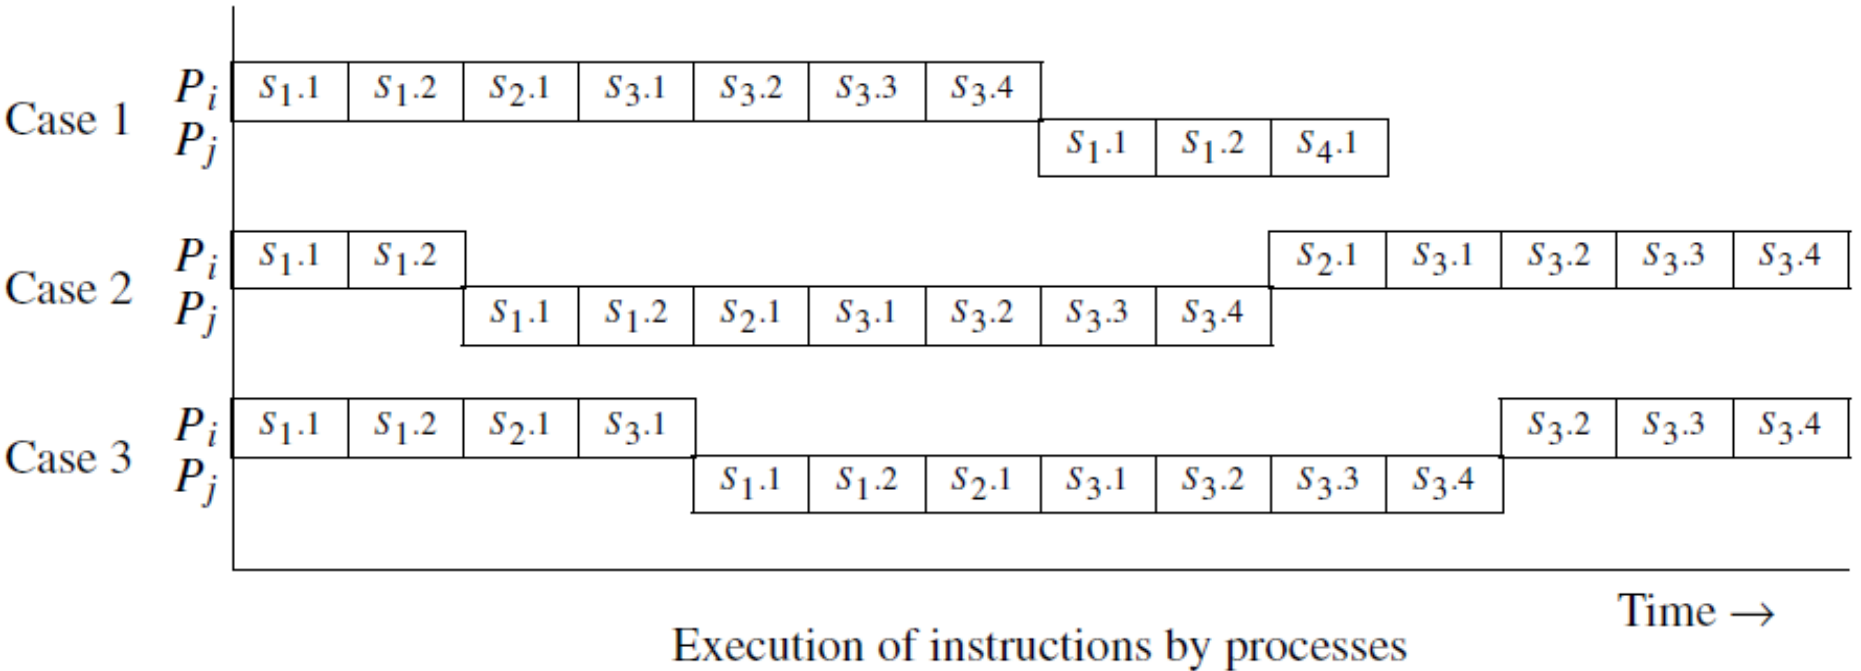
$y = y * 2$ ;

- 1 or 3 or 5 (non-deterministic)
- Race Condition: Thread A races against Thread B

# Race Condition


<u>Code of processes</u>		<u>Corresponding machine instructions</u>	
$S_1$	<b>if</b> $nextseatno \leq capacity$	$S_{1.1}$	Load $nextseatno$ in $reg_k$
		$S_{1.2}$	If $reg_k > capacity$ goto $S_4.1$
	<b>then</b>		
$S_2$	$allotedno := nextseatno;$	$S_{2.1}$	Move $nextseatno$ to $allotedno$
$S_3$	$nextseatno := nextseatno + 1;$	$S_{3.1}$	Load $nextseatno$ in $reg_j$
		$S_{3.2}$	Add 1 to $reg_j$
		$S_{3.3}$	Store $reg_j$ in $nextseatno$
		$S_{3.4}$	Go to $S_5.1$
	<b>else</b>		
$S_4$	$display$ “sorry, no seats available”	$S_{4.1}$	Display “sorry, . . .”
$S_5$	. . .	$S_{5.1}$	. . .

# Some execution cases



## More Definitions

- **Lock:** prevents someone from doing something
  - Lock before entering critical section and before accessing shared data
  - Unlock when leaving, after accessing shared data
  - Wait if locked
    - Important idea: all synchronization involves waiting
- Example: fix the milk problem by putting a lock on refrigerator
  - Lock it and take key if you are going to go buy milk



# Motivation: “Too much milk”

- Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

## Too Much Milk: Correctness

1. At most one person buys milk (not more than 1)
2. At least one person buys milk if needed



## Solution Attempt #1

- Leave a note
  - Place on fridge before buying
  - Remove after buying
  - Don't go to store if there's already a note
- Leaving/checking a note is atomic (word load/store)

```
if (noMilk) {  
    if (noNote) {  
        leave Note;  
        buy milk;  
        remove Note;  
    }  
}
```

# Attempt #1 in Action

A

```
if (noMilk) {  
    if (noNote) {  
  
        leave Note;  
        buy milk;  
        remove Note;  
    }  
}
```

B

```
if (noMilk) {  
    if (noNote) {  
  
        leave Note;  
        buy milk;  
        remove note;  
    }  
}
```

## Solution Attempt #2

*Leave Note;*

```
if (noMilk) {
```

```
  if (noNote) {
```

```
    leave Note;
```

```
    buy milk;
```

```
  }
```

```
}
```

```
remove Note;
```

But there's always a note –  
you just left one!

At least you don't  
buy milk twice...

# Solution Attempt #3

- Leave a named note – each person ignores their own

**A**

```
leave note A
if (noMilk) {
    if (noNote B) {
        buy milk
    }
}
remove note A;
```

**B**

```
leave note B
if (noMilk) {
    if (noNote A) {
        buy milk
    }
}
remove note B;
```

## Attempt #3 in Action

A

```
leave note A
if (noMilk) {
    if (noNote B) {
        buy milk
    }
}
```

remove note A

B

leave note B

```
if (noMilk) {
    if (noNote A) {
        buy milk
    }
    remove note B
```

## Solution Attempt #4

**A**

```
leave note A
while (note B) {
    do nothing
}
if (noMilk) {
    buy milk
}
remove note A;
```

**B**


```
leave note B
if (noNote A) {
    if (noMilk) {
        buy milk
    }
}
remove note B;
```

- This is a correct solution, but ...

# Issues with Solution 4

- Complexity
  - Proving that it works is hard
  - How do you add another thread?
- Busy-waiting
  - **A consumes CPU time to wait**
- Fairness
  - Who is more likely to buy milk?


# All Possible Executions" ?



a) One execution



b) Another execution



c) Another execution



# Relevant Definitions

- **Lock:** An object only one thread can hold at a time
  - **Provides** mutual exclusion
- Offers two **atomic** operations:
  - Lock.Acquire() – wait until lock is free; then grab
  - Lock.Release() – Unlock, wake up waiters

# Using Locks

```
MilkLock.Acquire()  
if (noMilk) {  
    buy milk  
}  
MilkLock.Release()
```

**But how do we implement this?**

**First, how do we use it?**

## Relevant Definitions

- **Lock:** An object only one thread can hold at a time
  - **Provides** mutual exclusion
- Offers two **atomic** operations:
  - Lock.Acquire() – wait until lock is free; then grab
  - Lock.Release() – Unlock, wake up waiters



# Mutexes

- Critical sections typically associated with mutual exclusion locks (*mutexes*)
- Only one thread can hold a given mutex at a time
- Acquire (lock) mutex on entry to critical section
  - Or block if another thread already holds it
- Release (unlock) mutex on exit
  - Allow one waiting thread (if any) to acquire & proceed

```
pthread_mutex_init(&m);  
pthread_mutex_lock(&m);    pthread_mutex_lock(&m);  
    hits = hits+1;          hits = hits+1;  
pthread_mutex_unlock(&m);  pthread_mutex_unlock(&m);  
  
    ↪ T1                    ↪ T2
```

# Using Locks

```
MilkLock.Acquire()  
if (noMilk) {  
    buy milk  
}  
MilkLock.Release()
```

But how do we implement this?

First, how do we use it?

# Implementing Too Much Milk with Locks

Thread A

Thread B

```
Lock.Acquire();  
if (noMilk){  
    buy milk;  
}  
Lock.Release();
```

```
Lock.Acquire();  
if (noMilk){  
    buy milk;  
}  
Lock.Release();
```

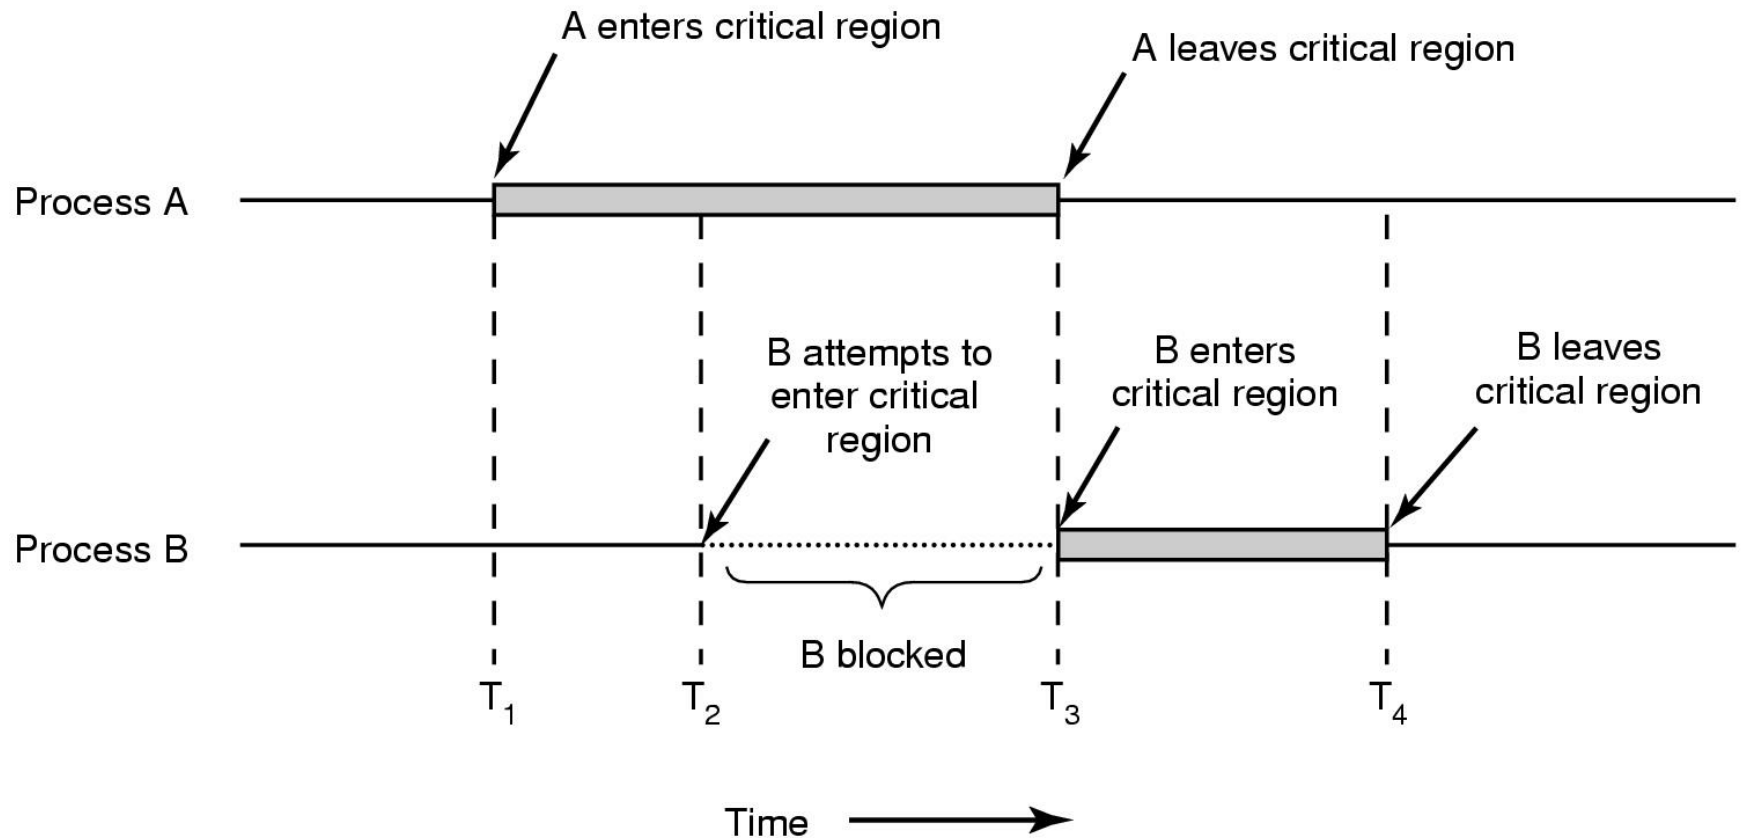
- This solution is clean and symmetric.
- How do we make Lock.Acquire and Lock.Release atomic?

# Support for synchronization?

Programs	Shared Programs
Higher-level API	Locks   Semaphores   Monitors   Send/Receive
Hardware	Load/Store   Disable Ints   Test&Set   Comp&Swap

- We are going to implement various higher-level synchronization primitives using atomic operations
  - Everything is pretty painful if only atomic primitives are load and store
  - Need to provide primitives useful at user-level

# Re-call Critical Section Problem



Mutual exclusion using critical regions



“Solution 1 ”:

**Disable interrupts** while  
holding lock

# Implementing Locks: Single Core

- Idea: A context switch can only happen (assuming threads don't yield) if there's an **interrupt**
- “Solution”: **Disable interrupts** while holding lock
- x86 has `cli` and `sti` instructions that only operate in system mode (PL=0)
  - Interrupts enabled bit in FLAGS register

# Interrupt Enable/Disable

```
Acquire() {  
    disable interrupts;  
}
```

```
Release() {  
    enable interrupts;  
}
```

- Problem: can stall the entire system

```
Lock.Acquire()  
While (1) {}
```

- Problem: What if we want to do I/O?

```
Lock.Acquire()  
Read from disk  
/* OS waits for (disabled) interrupt)!  
*/
```

# Implementation of Locks by Disabling Interrupts

- Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable

```
int value = FREE;
```



```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

```
Release() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    enable interrupts;  
}
```

## Atomic Read-Modify-Write instructions

- Problems with interrupt-based lock solution:
  - Can't give lock implementation to users
  - Doesn't work well on multiprocessor
    - Disabling interrupts on all processors requires messages and would be very time consuming
- Alternative: atomic instruction sequences
  - These instructions read a value from memory and write a new value atomically
  - Hardware is responsible for implementing this correctly
    - on both uniprocessors (not too hard)
    - and multiprocessors (requires help from cache coherence protocol)
  - Unlike disabling interrupts, can be used on both uniprocessors and multiprocessors

# An atomic test\_and\_set instruction

- Assume we add the following instruction
- **Test\_And\_Set Rx memory-location**
- Perform the following atomically:
  - Read the current value of memory-location into some processor register (Rx)
  - Set the memory-location to a 1

# Test\_and\_Set

```
boolean Test_and_Set( boolean memory[m] )
{ [
    if( memory[m] )    // lock denied
        return True;
    else {              // lock granted
        memory[m] = True;
        return False;
    }
}
]
```

# An atomic test\_and\_set instruction

## If the mutex is unlocked (=0)

- test\_and\_set will return 0 which will mean you have the mutex lock
- It will also set the mutex variable to 1

## if the mutex is locked (=1)

- test\_and\_set will return 1 which will mean you don't have the mutex lock
- It will also set the mutex variable to 1 (which is what it was anyway)

(mutual exclusion- “**mutex**”)



# Examples of Read-Modify-Write

- `test&set (&address) { /* most architectures */  
 result = M[address];  
 M[address] = 1;  
 return result;  
}`
- `swap (&address, register) { /* x86 */  
 temp = M[address];  
 M[address] = register;  
 register = temp;  
}`
- `compare&swap (&address, reg1, reg2) { /* 68000 */  
 if (reg1 == M[address]) {  
 M[address] = reg2;  
 return success;  
 } else {  
 return failure;  
 }  
}`

# Implementing Locks with test&set

- Simple solution:

```
int value = 0; // Free
Acquire() {
    while (test&set(value)); // while busy
}
Release() {
    value = 0;
}
```

```
test&set (&address) {
    result = M[address];
    M[address] = 1;
    return result;
}
```

- Simple explanation:

- If lock is free, test&set reads 0 and sets value=1, so lock is now busy. It returns 0 so while exits
- If lock is busy, test&set reads 1 and sets value=1 (no change). It returns 1, so while loop continues
- When we set value = 0, someone else can get lock

## Problem: Busy-Waiting for Lock



- Positives for this solution
  - Machine can receive interrupts
  - User code can use this lock
  - Works on a multiprocessor
- Negatives
  - Inefficient: busy-waiting thread will consume cycles waiting
  - Waiting thread may take cycles away from thread holding lock!
  - **Priority Inversion**: If busy-waiting thread has higher priority than thread holding lock  $\Rightarrow$  no progress!
- Priority Inversion problem with original Mars Rover

# Locks using test&set vs. Interrupts

- Compare to “disable interrupt” solution

```
int value = FREE;
```



```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
        // Enable interrupts?  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

```
Release() {  
    disable interrupts;  
    if (anyone on wait queue) {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    enable interrupts;  
}
```

✎ Basically replace

- disable interrupts → while (test&set(guard));
- enable interrupts → guard = 0;

# Better Locks using test&set

- Can we build test&set locks without busy-waiting?
  - Can't entirely, but can minimize!
  - Idea: only busy-wait to atomically check lock value

```
int guard = 0;  
int value = FREE;
```

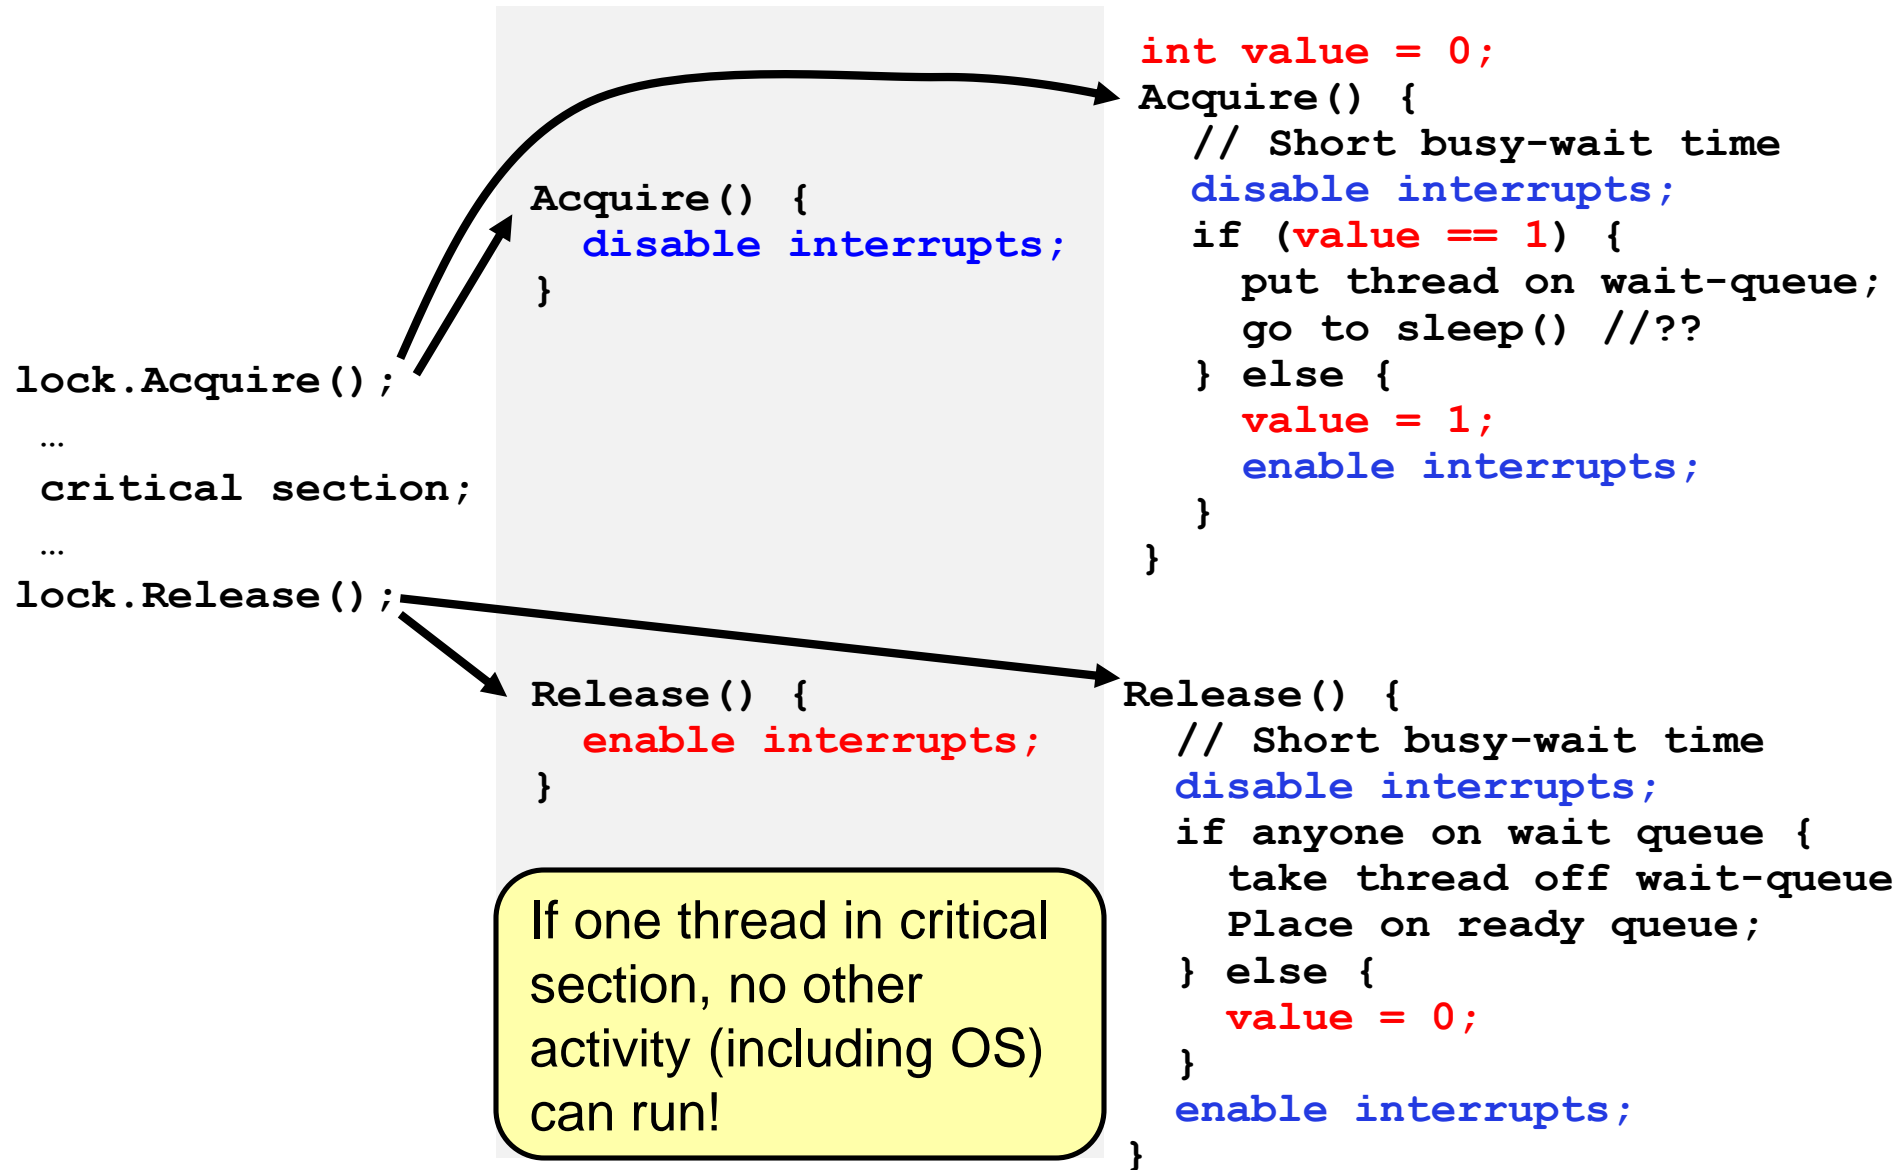


```
Acquire() {  
    // Short busy-wait time  
    while (test&set(guard));  
    if (value == BUSY) {  
        put thread on wait queue;  
        go to sleep() & guard = 0;  
    } else {  
        value = BUSY;  
        guard = 0;  
    }  
}
```

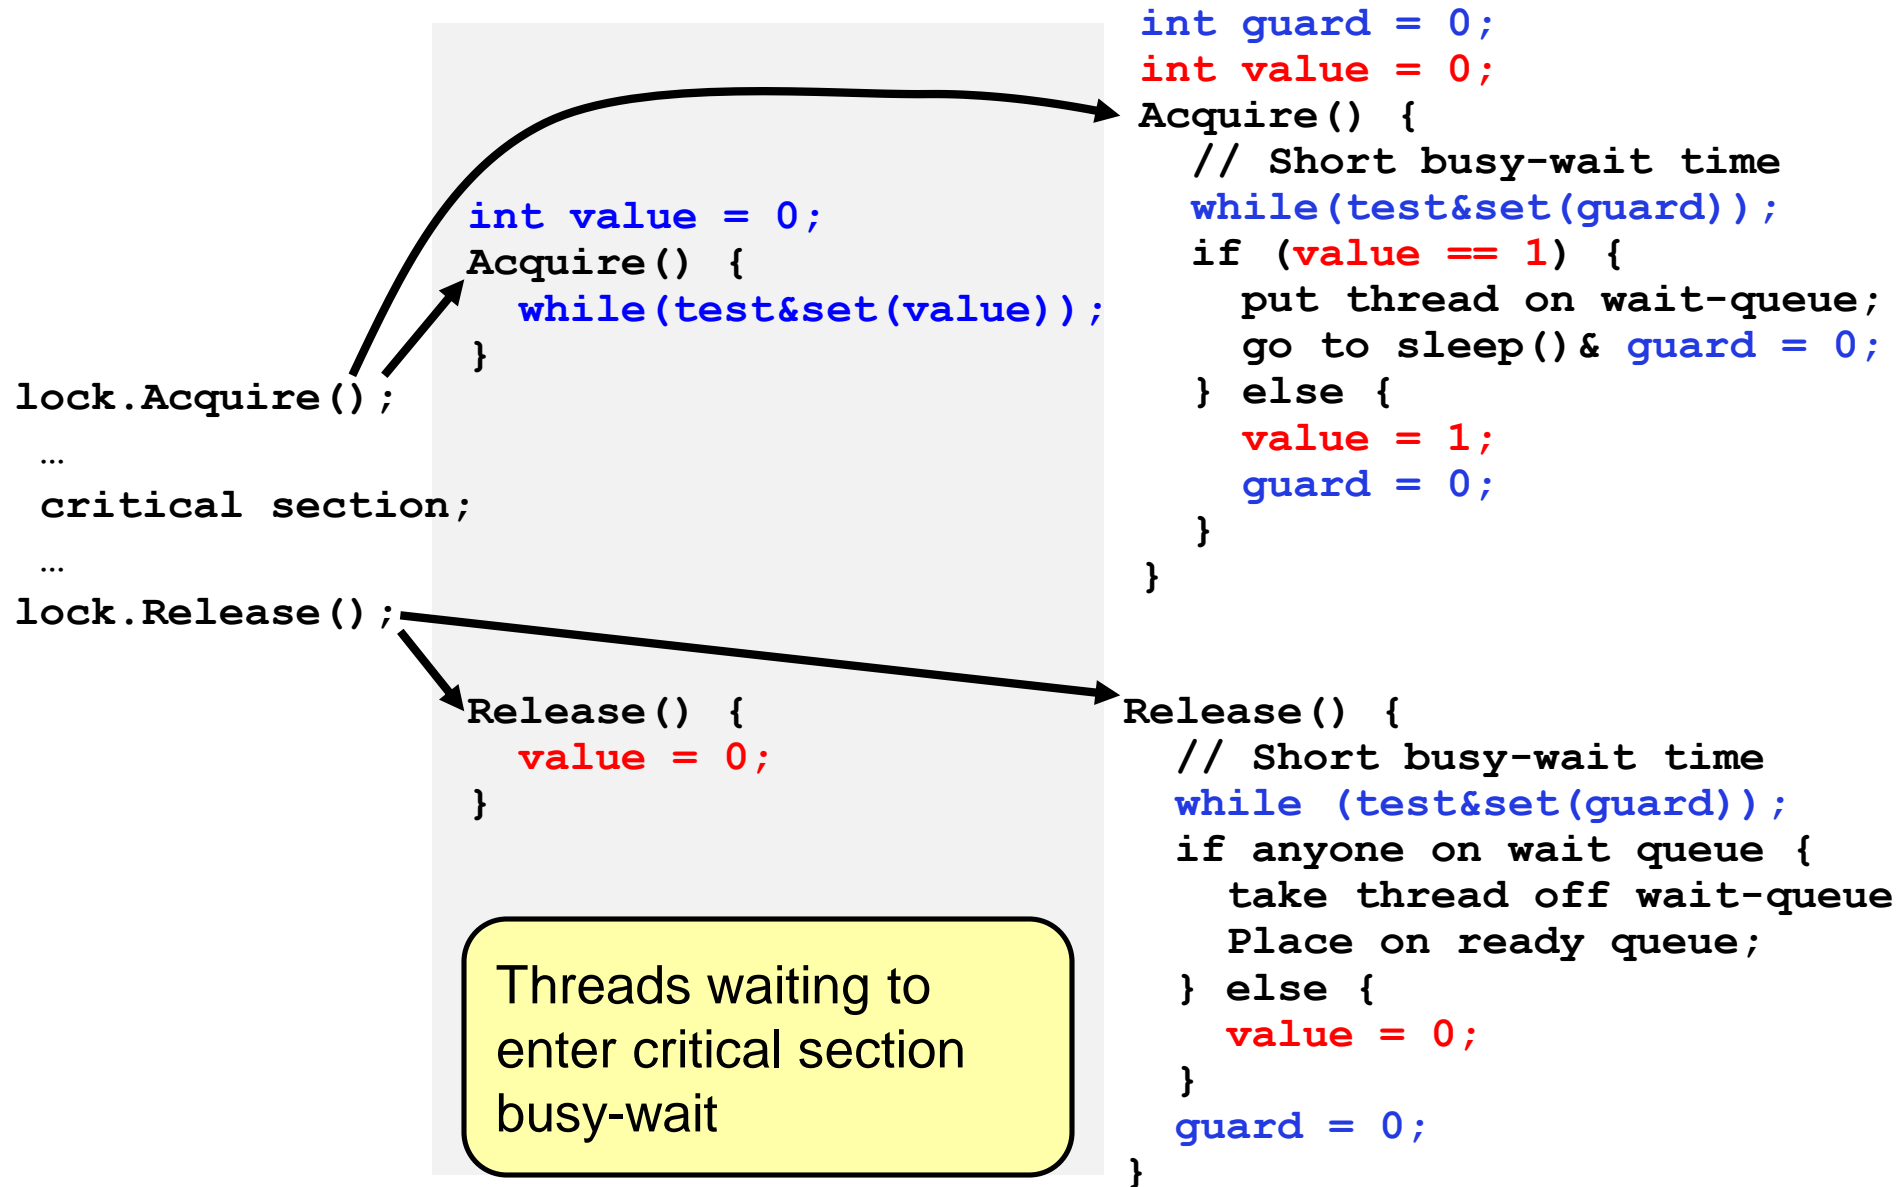
```
Release() {  
    // Short busy-wait time  
    while (test&set(guard));  
    if anyone on wait queue {  
        take thread off wait queue  
        Place on ready queue;  
    } else {  
        value = FREE;  
    }  
    guard = 0;  
}
```

- Note: sleep has to be sure to reset the guard variable
  - Why can't we do it just before or just after the sleep?

# Recap: Locks



# Recap: Locks



# Support for synchronization?

Programs	Shared Programs
Higher-level API	Locks   Semaphores   Monitors   Send/Receive
Hardware	Load/Store   Disable Ints   Test&Set   Comp&Swap

- We are going to implement various higher-level synchronization primitives using atomic operations
  - Everything is pretty painful if only atomic primitives are load and store
  - Need to provide primitives useful at user-level



# Peterson's solution

Peterson's algorithm (aka Peterson's solution) is a concurrent programming algorithm for **mutual exclusion** that allows **two tasks** to **share** a single use **resource** without conflict, **using only shared memory** for communication. It was formulated by Gary L. Peterson in 1981.

# Faulty Algorithm 1 - Turn taking

- The shared variable **turn** is initialized (to 0 or 1) before executing any  $P_i$
- $P_i$ 's critical section is executed iff  $\text{turn} = i$
- $P_i$  is **busy waiting** if  $P_j$  is in CS

```
Process  $P_i$ :  
    //  $i, j = 0$  or  $1$   
repeat  
    while ( $\text{turn} \neq i$ ) {};  
    CS  
     $\text{turn} := j$ ;  
    RS  
forever
```

Process P0:  
repeat

while (turn!=0) {};

CS

turn:=1;

RS

forever

Process P1:  
repeat


while (turn!=1) {};

CS

turn:=0;

RS

forever



Faulty Algorithm 1 side-by-side view

# Analysis

- Achieves Mutual Exclusion (busy wait)
- **But** Progress requirement is not satisfied since it requires strict alternation of CS's.
  - If one process requires its CS more often than the other, it can't get it.

# Faulty Algorithm 2 – Ready flag

- Keep a Boolean variable for each process: flag[0] and flag[1]
- $P_i$  signals that it is ready to enter its CS by: flag[i]:=true but waits until the other has finished its CS.

```
Process  $P_i$ :  
repeat  
    flag[i] := true;  
    while(flag[j]) {};  
    CS  
    flag[i] := false;  
    RS  
forever
```

Process P0:

repeat

flag[0] := true;

while(flag[1]){};

CS

flag[0] := false;

RS

forever

Process P1:

repeat

flag[1] := true;


while(flag[0]){};

CS

flag[1] := false;

RS

forever



Faulty Algorithm 2 side-by-side view

# Analysis

- Mutual Exclusion is satisfied but not the progress requirement
- For the (interleaved) sequence:
  - `flag[0]:=true`
  - `flag[1]:=true`
- Both processes will wait forever

# Algorithm 3

## (Peterson's Algorithm)

- Initialization:  
flag[0]:=flag[1]:=false  
turn:= 0 or 1
- Wish to enter CS specified  
by flag[i]:=true
- Even if both flags go up,  
and no matter how the  
instructions are  
interleaved,
  - ..turn will always end up  
as either 0 or 1

Process P<sub>i</sub>:

repeat

flag[i]:=true;

// I want in

turn:=j;

// but you can go first!

while(flag[j]&& turn==j) ;

CS

flag[i]:=false;

// I'm done

RS

forever



# Peterson's Algorithm

Process P0:

repeat

flag[0]:=true;

// 0 wants in

turn:= 1;

// 0 gives a chance to 1

while

(flag[1]&turn=1) ;

CS

flag[0]:=false;

// 0 is done

RS

forever

Process P1:

repeat

flag[1]:=true;

// 1 wants in

turn:=0;

// 1 gives a chance to 0

while

(flag[0]&turn=0) ;

CS

flag[1]:=false;

// 1 is done

RS

forever


Peterson's algorithm side-by-side view

# Peterson's Algorithm: Proof of Correctness

- Mutual exclusion holds since:
  - For both  $P_0$  and  $P_1$  to be in their CS
    - both  $\text{flag}[0]$  and  $\text{flag}[1]$  must be true and:
    - $\text{turn}=0$  and  $\text{turn}=1$  (at same time): impossible



# Semaphores



Edsger Dijkstra  
1930 - 2002

- Semaphores are a kind of generalized lock
  - First defined by Dijkstra in late 60s
  - Main synchronization primitive used in original UNIX (& Pintos)
- Definition: a Semaphore has a non-negative integer value and supports the following two operations:
  - **P() or down():** atomic operation that waits for semaphore to become positive, then decrements it by 1
  - **V() or up():** an atomic operation that increments the semaphore by 1, waking up a waiting P, if any

**P()** stands for “*proberen*” (to test) and **V()** stands for “*verhogen*” (to increment) in Dutch

# Semaphores: Key Concepts

- Like locks, a semaphore supports two atomic operations, Semaphore.Wait() and Semaphore.Signal().

```
S.Wait()           // wait until semaphore S
                    // is available
```

```
<critical section>
```

```
S.Signal()         // signal to other processes
                    // that semaphore S is free
```

- Each semaphore supports a queue of processes that are waiting to access the critical section (e.g., to buy milk).
- If a process executes **S.Wait()** and semaphore S is free (non-zero), it continues executing. If semaphore S is not free, the OS puts the process on the wait queue for semaphore S.
- A **S.Signal()** unblocks one process on semaphore S's wait queue.

# Semaphore

- Does not require **busy-waiting**
  - CPU is not held unnecessarily while the process is waiting
- A Semaphore  $S$  is
  - A data structure with an integer variable  $S.value$  and a queue  $S.list$  of processes (shared variable)
  - The data structure can only be accessed by two **atomic** operations,  $wait(S)$  and  $signal(S)$  (also called  $down(S)$ ,  $P(S)$  and  $Up(s)$ ,  $V(S)$ )
- Value of the semaphore  $S$  = value of the integer  $S.value$

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore
```

# Semaphore

## Wait(S)      $S \leq 0$ semaphore variable

- When a process P executes the wait(S) and finds
- $S = 0$ 
  - Process must wait  $\Rightarrow$  block()
  - Places the process into a waiting queue associated with S
  - Switch from running to waiting state

## Signal(S)

When a process P executes the signal(S)

- Check, if some other process Q is waiting on the semaphore S
- Wakeup(Q)
- Wakeup(Q) changes the process from waiting to ready state

# Semaphore (wait and signal)

- Implementation of wait:

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

List of PCB



Atomic/  
Indivisible

- Implementation of signal:

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```

**Note:** which process is picked for unblocking may depend on policy.

# Binary Semaphores: Example

- Too Much Milk using locks:

Thread A

```
Lock.Acquire();  
if (noMilk){  
    buy milk;  
}  
Lock.Release();
```

Thread B

```
Lock.Acquire();  
if (noMilk){  
    buy milk;  
}  
Lock.Release();
```

- Too Much Milk using semaphores:

Thread A

```
Semaphore.Wait();  
if (noMilk){  
    buy milk;  
}  
Semaphore.Signal();
```

Thread B

```
Semaphore.Wait();  
if (noMilk){  
    buy milk;  
}  
Semaphore.Signal();
```



Semaphore -> P() (*Passeren*; wait)

    If  $\text{sem} > 0$ , then decrement sem by 1

    Otherwise “wait” until  $\text{sem} > 0$  and then  
    decrement


Semaphore -> V() (*Vrijgeven*; signal)

    Increment sem by 1

    Wake up a thread waiting in P()

# Semaphores Like Integers Except

- Semaphores are like integers, except
  - No negative values
  - Only operations allowed are P and V – can't read or write value, except to set it initially
  - Operations must be atomic
    - Two P's together can't decrement value below zero
    - Similarly, thread going to sleep in P won't miss wakeup from V – even if they both happen at same time
- Semaphore from railway analogy
  - Here is a semaphore initialized to 2 for resource control:



# Signal and Wait: Example

P1: S.Wait();

S.Wait();

S.Signal();

S.Signal();

P2: S.Wait();

S.Signal();

P1: S->Wait();  
P2: S->Wait();  
P1: S->Wait();  
P2: S->Signal();  
P1: S->Signal();  
P1: S->Signal();

value	process state: execute or block		
	Queue	P1	P2
2	empty	execute	execute
1	empty	execute	execute
0	empty	execute	execute
-1	P1	blocked	execute
0	empty	execute	execute
1	empty	execute	execute
2	empty	execute	execute

# Two Uses of Semaphores

- Mutual Exclusion (initial value = 1)
  - Also called “Binary Semaphore”.
  - Can be used for mutual exclusion:

```
semaphore.P();  
// Critical section goes here  
semaphore.V();
```
- Scheduling Constraints (initial value = 0)
  - Allow thread 1 to wait for a signal from thread 2, i.e., thread 2 **schedules** thread 1 when a given **constrained** is satisfied
  - Example: suppose you had to implement ThreadJoin which must wait for thread to terminate:

Initial value of semaphore = 0

```
ThreadJoin {  
    semaphore.P();  
}
```

```
ThreadFinish {  
    semaphore.V();  
}
```



# Deadlocks

**Deadlock:** A condition where two or more threads are waiting for an event that can only be generated by these same threads.

Example:

Process A:

```
printer.Wait();  
disk.Wait();
```

```
// copy from disk  
// to printer
```

```
printer.Signal();  
disk.Signal();
```

Process B:

```
disk.Wait();  
printer.Wait();
```

```
// copy from disk  
// to printer
```

```
printer.Signal();  
disk.Signal();
```

# Deadlock and Starvation

- Let  $S$  and  $Q$  be two semaphores initialized to 1

$P_0$

wait (S);

wait (Q);

.

.

.

signal (S);

signal (Q);

$P_1$

wait (Q);

wait (S);

.

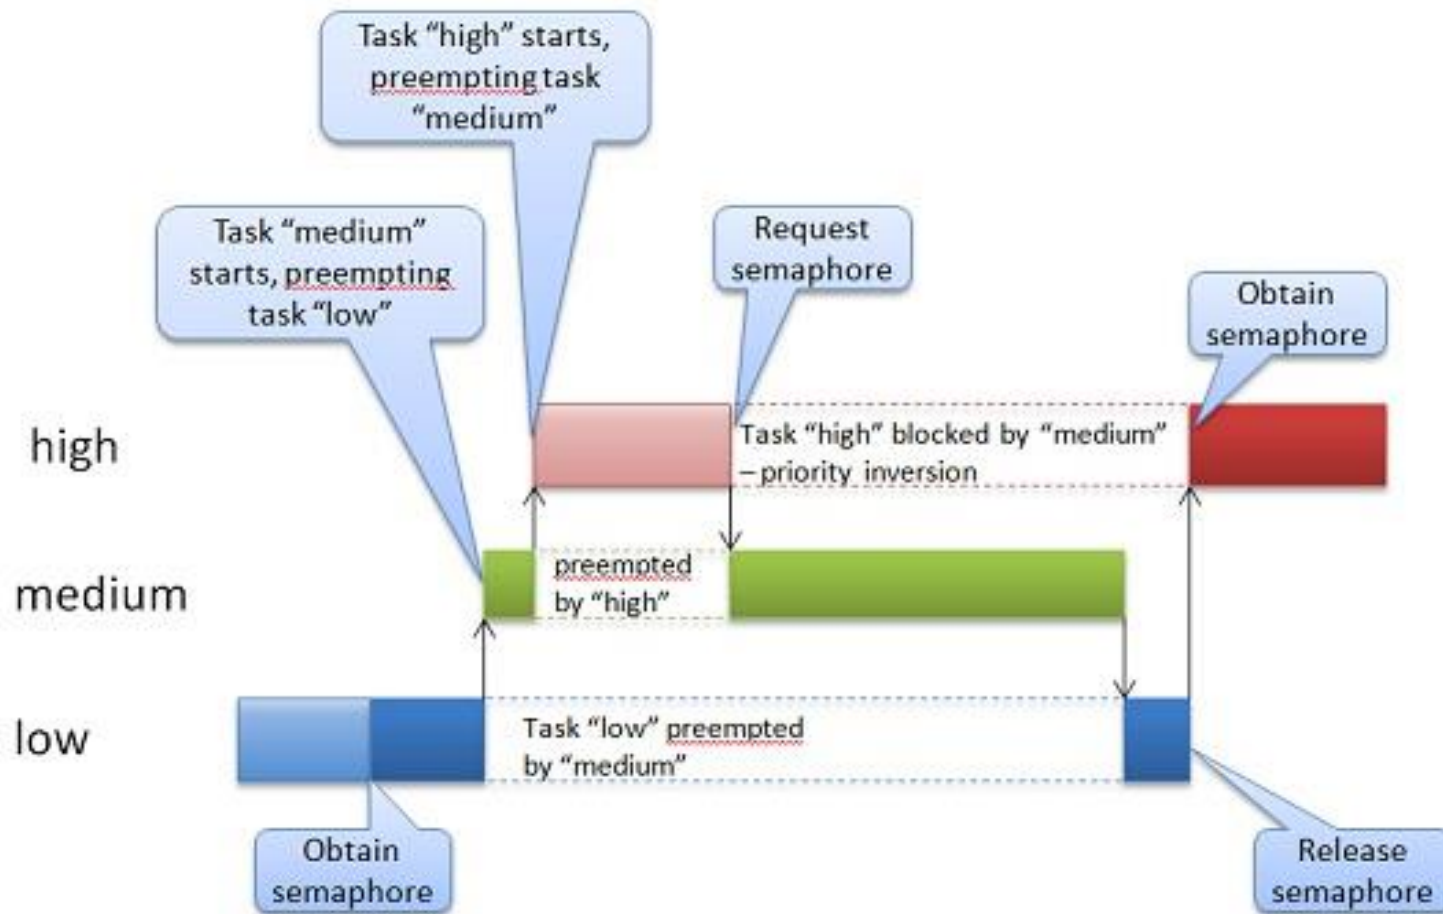
.

.

signal (Q);

signal (S);

- Starvation** – indefinite blocking
  - LIFO queue
  - A process may never be removed from the semaphore queue in which it is suspended



<https://www.rapitasystems.com/blog/what-really-happened-to-the-software-on-the-mars-pathfinder-spacecraft>

# Ordering Execution of Processes using Semaphores

- Execute statement  $B$  in  $P_j$  only after statement  $A$  executed in  $P_i$
- Use semaphore  $flag$  initialized to 0
- Code:

$P_i$	$P_j$
$\vdots$	$\vdots$
Stmt. $A$	$wait(flag)$
$signal(flag)$	Stmt. $B$

- Multiple such points of synchronization can be enforced using one or more semaphores



# Synchronization Problems- Real-World Examples

- **Producer-consumer**

- Audio/Video player: network and display threads; shared buffer

- Web servers: master thread and slave thread

- **Reader-writer**

- Banking system: read account balances versus update

- **Dining Philosophers**

- Cooperating processes that need to share limited resources


- Set of processes that need to lock multiple resources

- Disk and tape (backup),


- Travel reservation: hotel, airline, car rental databases

# Web Server: Thread Pools

- **Bounded** pool of worker threads
  - Allocated in **advance**: no thread creation overhead
  - **Queue** of pending requests
  - **Limited number** of requests in progress




# Readers/Writers Problem



- Motivation: Consider a shared database
  - Two classes of users:
    - Readers – never modify database
    - Writers – read and modify database
  - Is using a single lock on the whole database sufficient?
    - Like to have many readers at the same time
    - Only one writer at a time

# Basic Readers/Writers Solution



- Correctness Constraints:
  - Readers can access database when no writers
  - Writers can access database when no readers or writers
  - Only one thread manipulates state variables at a time
- Basic structure of a solution:
  - **Reader()**
    - Wait until no writers
    - Access database
    - Check out – wake up a waiting writer
  - **Writer()**
    - Wait until no active readers or writers
    - Access database
    - Check out – wake up waiting readers or writer
  - State variables (Protected by a lock called “lock”):
    - int AR: Number of active readers; initially = 0
    - int WR: Number of waiting readers; initially = 0
    - int AW: Number of active writers; initially = 0
    - int WW: Number of waiting writers; initially = 0
    - Condition okToRead = NIL
    - Condition okToWrite = NIL

# Solution to Bounded Buffer (coke machine)

```
Semaphore fullSlots = 0;    // Initially, no coke
Semaphore emptySlots = bufSize;
                             // Initially, num empty slots
Semaphore mutex = 1;        // No one using machine
```

```
Producer(item) {
    semaP(&emptySlots);    // Wait until space
    semaP(&mutex);          // Wait until machine free
    Enqueue(item);
    semaV(&mutex);
    semaV(&fullSlots);     // Tell consumers there is
                           // more coke
}
```

```
Consumer() {
    semaP(&fullSlots);     // Check if there's a coke
    semaP(&mutex);          // Wait until machine free
    item = Dequeue();
    semaV(&mutex);
    semaV(&emptySlots);    // tell producer need more
    return item;
}
```


emptySlots  
signals space

fullSlots signals coke

Critical sections  
using mutex  
protect integrity  
of the queue



# Dining Philosophers Problem



- Five philosophers, each either eats or thinks
- Share a circular table with five chopsticks
- Thinking: do nothing
- Eating => need two chopsticks, try to pick up two closest chopsticks
  - Block if neighbor has already picked up a chopstick
- After eating, put down both chopsticks and go back to thinking

# Dining Philosophers attempt 1

```
Semaphore  chopsticks[5];

do{
    wait(chopstick[i]);  // left chopstick
    wait(chopstick[(i+1)%5 ]); // right chopstick
    // eat
    signal(chopstick[i]);  // left chopstick
    signal(chopstick[(i+1)%5 ]); // right chopstick
    // think
} while(TRUE);
```

# Dining Philosophers attempt 2

```
monitor DP {
    enum { THINKING; HUNGRY,
EATING) state [5] ;
    condition self [5];

void synchronized pickup (int i) {
    state[i] = HUNGRY;
    test(i);
    if (state[i] != EATING)
        self[i].wait;
}


void synchronized putdown (int i) {
    state[i] = THINKING;
    //test left and right neighbors
    test((i + 4) % 5);
    test((i + 1) % 5);
}

void test (int i) {
    if ( (state[(i + 4) % 5] != EATING)&&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {
        state[i] = EATING ;
        self[i].signal () ;
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
}
```



# Producer-consumer with a bounded buffer



- Problem Definition
  - Producer puts things into a shared buffer
  - Consumer takes them out
  - Need synchronization to coordinate producer/consumer
- Don't want producer and consumer to have to work in lockstep, so put a fixed-size buffer between them
  - Need to synchronize access to this buffer
  - Producer needs to wait if buffer is full
  - Consumer needs to wait if buffer is empty
- Example: Coke machine
  - Producer can put limited number of cokes in machine
  - Consumer can't take cokes out if machine is empty



# Full Solution to Bounded Buffer

```
Semaphore fullSlots = 0;  
Semaphore emptySlots = bufSize;  
  
Semaphore mutex = 1;
```

// Initially, no coke

// Initially, num empty slots

// No one using machine

```
Producer(item) {  
    emptySlots.P();  
    mutex.P();  
    Enqueue(item);  
    mutex.V();  
    fullSlots.V();  
}
```

// Wait until space

// Wait until machine free

// Tell consumers there is  
// more coke

```
Consumer() {  
    fullSlots.P();  
    mutex.P();  
    item = Dequeue();  
    mutex.V();  
    emptySlots.V();  
    return item;  
}
```

// Check if there's a coke

// Wait until machine free

// tell producer need more

## Motivation for Monitors and Condition Variables

- Cleaner idea: Use *locks* for mutual exclusion and *condition variables* for scheduling constraints
- **Monitor**: a lock and zero or more condition variables for managing concurrent access to shared data
  - Some languages like Java provide this natively
  - Most others use actual locks and condition variables

# Motivation for Monitors and Condition Variables

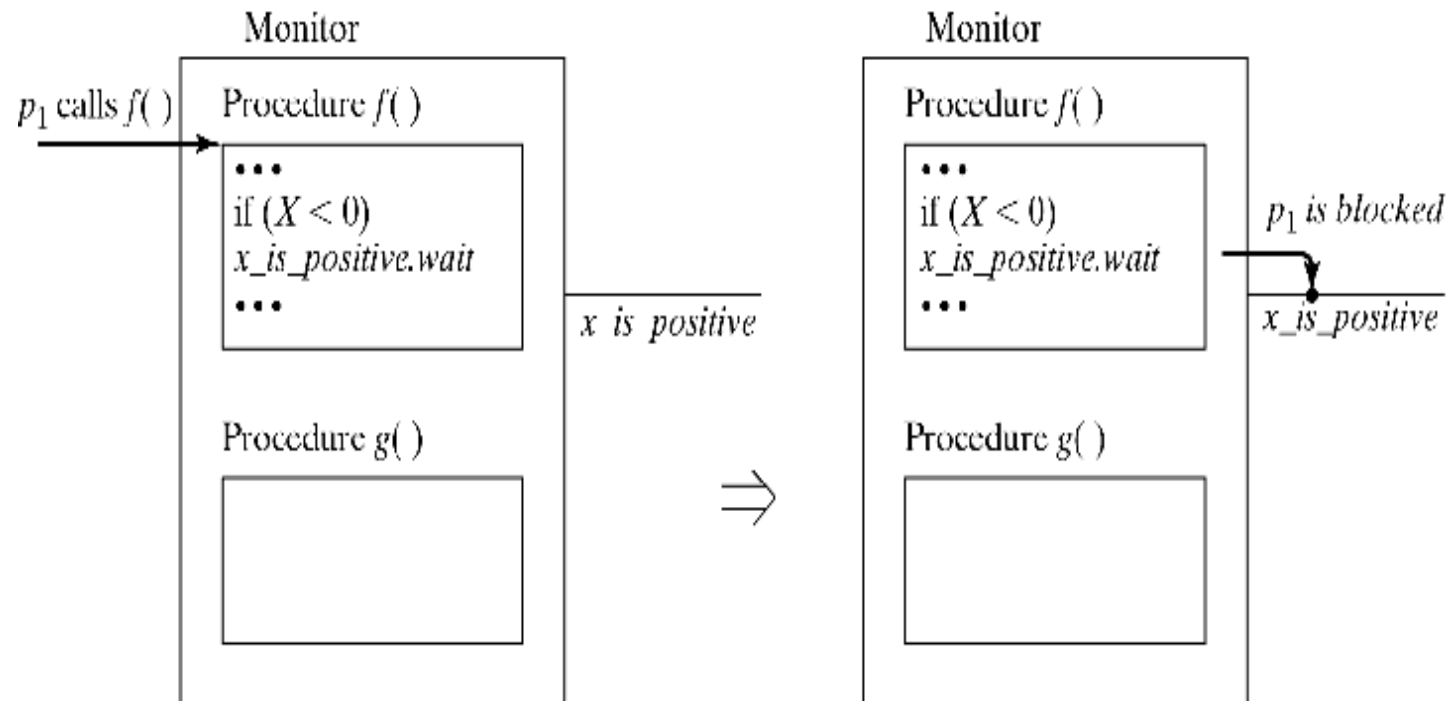
- Semaphores are a huge step up; just think of trying to do the bounded buffer with only loads and stores
- Problem is that semaphores are dual purpose:
  - They are used for both mutex and scheduling constraints
  - Example: the fact that flipping of P's in bounded buffer gives deadlock is not immediately obvious. How do you prove correctness to someone?

# Motivation for Monitors and Condition Variables

- Cleaner idea: Use *locks* for mutual exclusion and *condition variables* for scheduling constraints
- **Monitor**: a lock and zero or more condition variables for managing concurrent access to shared data
  - Some languages like Java provide this natively
  - Most others use actual locks and condition variables

# Monitors

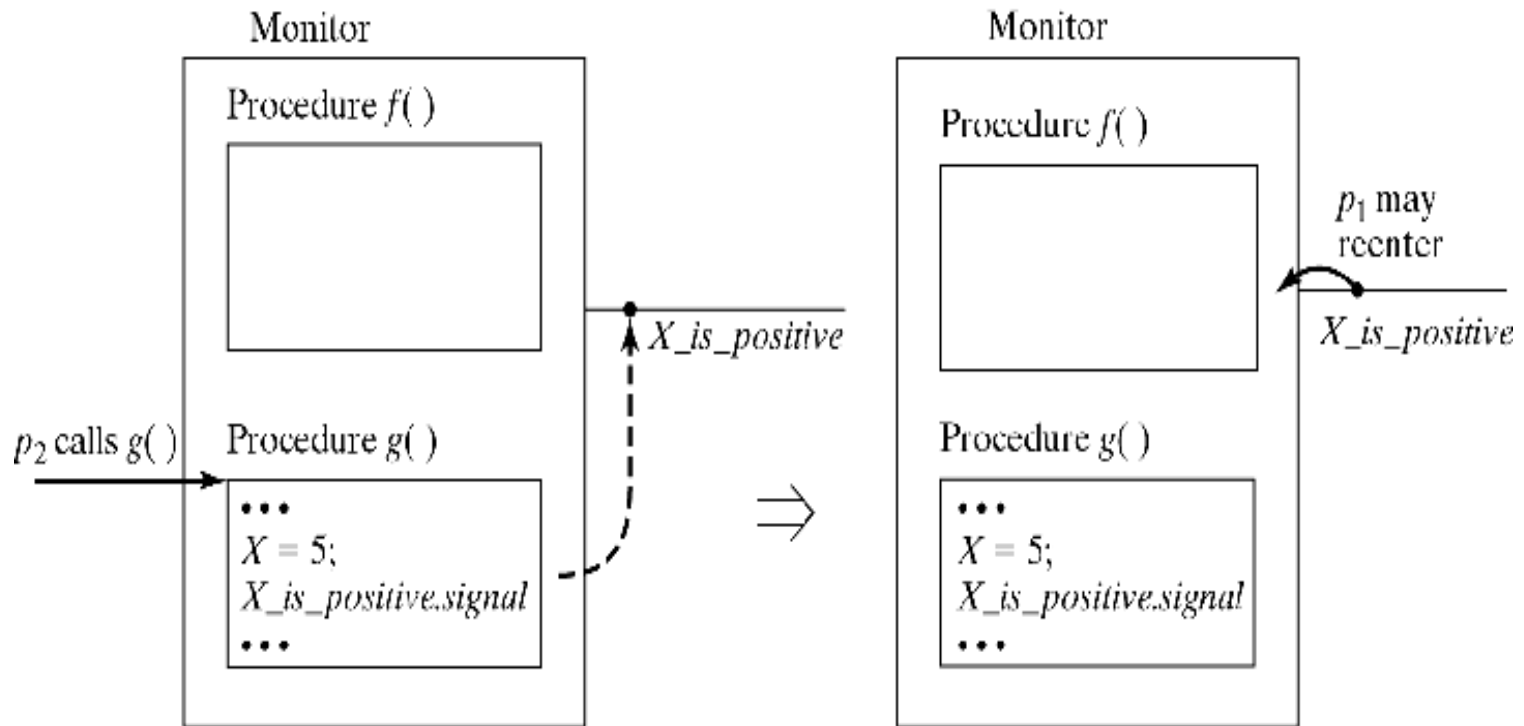
- Example: process p1 needs to wait until a variable X becomes positive before it can proceed



- p1 steps out to wait on queue associated with condition variable *X\_is\_positive*

# Monitors

- Another process may execute `X_is_positive.signal` to wake up `p1` when `X` becomes positive



- process `p1` may reenter the monitor, however ...

# Monitors

- Design Issue:
    - After c.signal, there are **2 ready processes**:
      - The **calling** process that did the c.signal
      - The **waiting** process that the c.signal woke up
    - Which should continue?  
(Only one can be executing inside the monitor!)
- Two different approaches
- **Hoare** monitors
  - **Mesa-style** monitors

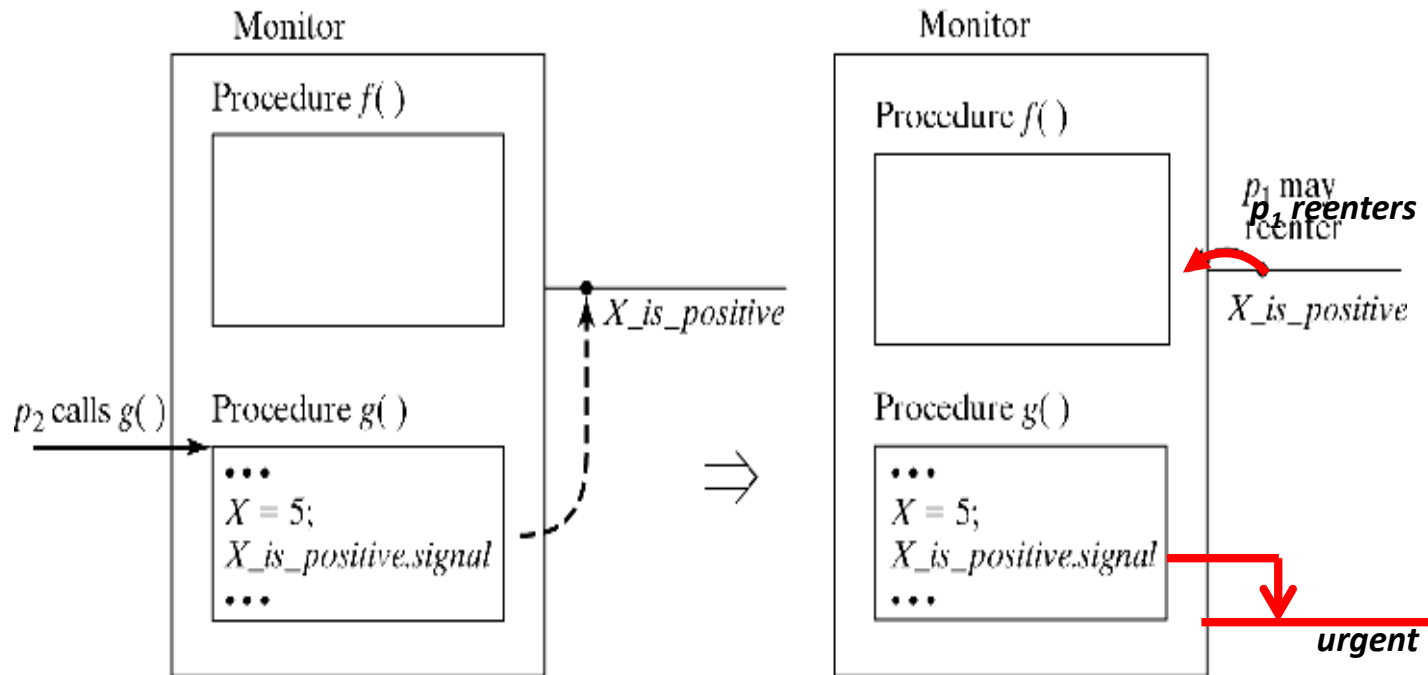


# Hoare Monitors

- Introduced by Tony Hoare in a 1974  
[http://wikipedia.org/wiki/C. A. R. Hoare](http://wikipedia.org/wiki/C._A._R._Hoare)
- First implemented by Per Brinch Hansen in Concurrent Pascal  
[http://wikipedia.org/wiki/Per Brinch Hansen](http://wikipedia.org/wiki/Per_Brinch_Hansen)
- Approach taken by Hoare monitor:
  - After c.signal,
    - **Awakened process continues**
    - **Calling process is suspended**, and placed on **high-priority** queue

# Hoare Monitors

- Effect of c.signal

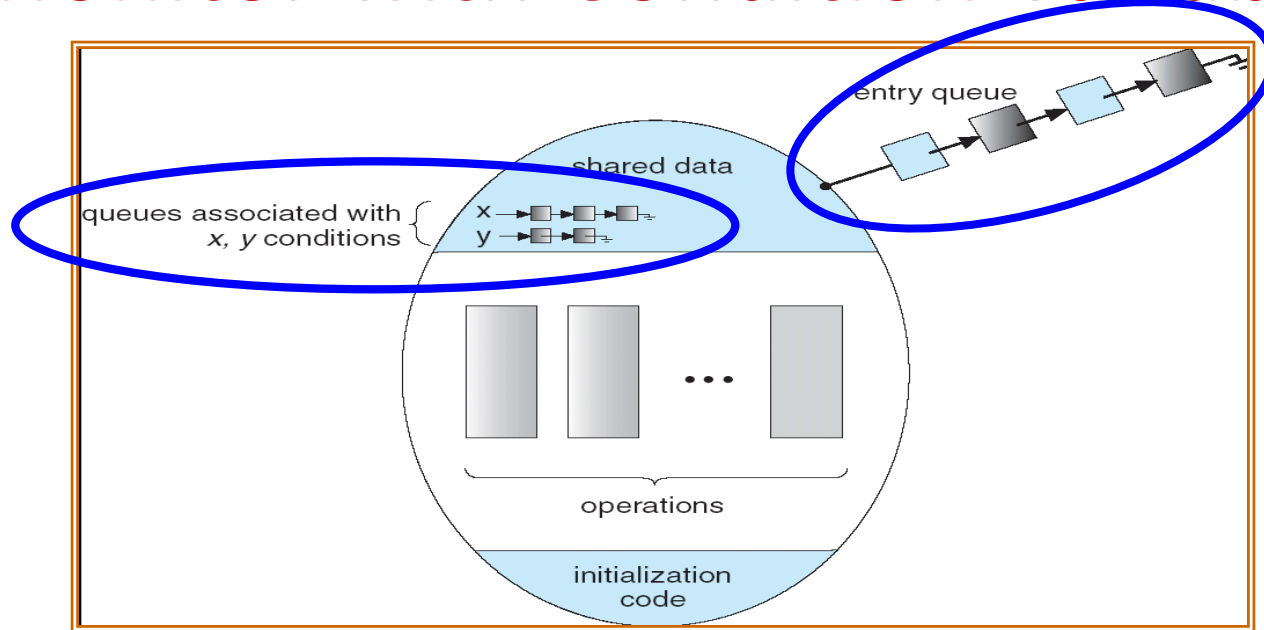


- Signaling process has the highest priority
  - It reenters as soon as p1 terminates

# Mesa-style monitors

- After issuing `c.signal`, **calling process continues** executing
- Problem: **Condition cannot be guaranteed** after wakeup
  - Assume that  $p_1$  and  $p_2$  are waiting for some condition  $c$
  - Caller could satisfy  $c$  and wake up both processes
  - Assume  $p_1$  starts and makes the condition false again
  - When  $p_2$  resumes,  $c$  is not guaranteed to be true
- Solution: process must **retest**  $c$  after wakeup
  - instead of:    `if (!condition) c.wait`
  - use:            `while (!condition) c.wait`
- This form of signal is sometimes called **notify**

# Monitor with Condition Variables



- **Lock:** the lock provides mutual exclusion to shared data
  - Always acquire before accessing shared data structure
  - Always release after finishing with shared data
  - Lock initially free
- **Condition Variable:** a queue of threads waiting for something *inside* a critical section
  - Key idea: make it possible to go to sleep inside critical section by atomically releasing lock at time we go to sleep

# Simple Monitor Example

- Here is an (infinite) synchronized queue

```
Lock lock;  
Queue queue;
```

```
AddToQueue(item) {  
    lock.Acquire();           // Lock shared data  
    queue.enqueue(item);      // Add item  
    lock.Release();           // Release Lock  
}
```

```
RemoveFromQueue() {  
    lock.Acquire();           // Lock shared data  
    item = queue.dequeue();    // Get next item or null  
    lock.Release();           // Release Lock  
    return(item);              // Might return null  
}
```

- Not very interesting use of “Monitor”
  - It only uses a lock with no condition variables
  - Cannot put consumer to sleep if no work!

# Complete Monitor Example (with condition variable)

- Here is an (infinite) synchronized queue

```
Lock lock;
```

```
Condition dataready;
```

```
Queue queue;
```

```
AddToQueue(item) {
```

```
    lock.Acquire();
```

```
    queue.enqueue(item);
```

```
    dataready.signal();
```

```
    lock.Release();
```

```
}
```

```
// Get Lock
```

```
// Add item
```

```
// Signal any waiters
```

```
// Release Lock
```

```
RemoveFromQueue() {
```

```
    lock.Acquire();
```

```
    while (queue.isEmpty()) {
```

```
        dataready.wait(&lock); // If nothing, sleep
```

```
    }
```

```
    item = queue.dequeue();
```

```
    lock.Release();
```


```
    return(item);
```

```
}
```

```
// Get Lock
```

```
// Get next item
```

```
// Release Lock
```



# Summary

- Locks construction based on atomic seq. of instructions
  - Must be very careful not to waste/tie up machine resources
    - Shouldn't spin wait for long
  - Key idea: Separate lock variable, use hardware mechanisms to protect modifications of that variable
- Semaphores
  - Generalized locks
  - Two operations: **P()**, **V()**
- Monitors: A synchronous object plus one or more condition variables
  - Always acquire lock before accessing shared data
  - Use condition variables to wait inside critical section
    - Three Operations: **Wait()**, **Signal()**, and **Broadcast()**