

Deadlock

Deadlock

- Let S and Q be two semaphores initialized to 1

P_0

```
wait (S);  
wait (Q);  
.  
.  
.  
signal (S);  
signal (Q);
```

P_1

```
wait (Q);  
wait (S);  
.  
.  
.  
signal (Q);  
signal (S);
```


Deadlocks

Deadlock can be defined formally as follows:

A set of processes is deadlocked if each process in the set is waiting for an event that only another process in the set can cause.

Types of Scheduling Problems

- Starvation – thread fails to make progress for an indefinite period of time
- Deadlock – starvation due to a *cycle of waiting* among a set of threads
 - each thread waits for some other thread in the cycle to take some action



Deadlock with Space

Thread A

AllocateOrWait(1 MB)

AllocateOrWait(1 MB)

Free(1 MB)

Free(1 MB)

Thread B

AllocateOrWait(1 MB)

AllocateOrWait(1 MB)

Free(1 MB)

Free(1 MB)

If only **2 MB of space**, we get same deadlock situation

Deadlock with Locks: “Lucky” Case

Thread A

`x.Acquire();`

`y.Acquire();`

`...`

`y.Release();`

`x.Release();`

Thread B

`y.Acquire();`

`x.Acquire();`

`...`


`x.Release();`

`y.Release();`


Sometimes schedule won't trigger deadlock

Resource-Allocation Graph


- Process




- Resource Type with 4 instances



- P_i requests instance of R_j




- P_i is holding an instance of R_j



Resources


- Resources – passive entities needed by threads to do their work
 - CPU time, disk space, memory
- Two types of resources:
 - Preemptable – can take it away
 - CPU, Embedded security chip
 - Non-preemptable – must leave it with the thread
 - Disk space, printer, chunk of virtual address space
 - Critical section
- Resources may require exclusive access or may be sharable
 - Read-only files are typically sharable
 - Printers are not sharable during time of printing
- One of the major tasks of an operating system is to manage resources



Resource-Allocation Graph

- System Model

- A set of Threads T_1, T_2, \dots, T_n
- Resource types R_1, R_2, \dots, R_m
CPU cycles, memory space, I/O devices
- Each resource type R_i has W_i instances
- Each thread utilizes a resource as follows:
 - `Request()` / `Use()` / `Release()`




- Resource-Allocation Graph:

- V is partitioned into two types:
 - $T = \{T_1, T_2, \dots, T_n\}$, the set threads in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$, the set of resource types in system
- request edge – directed edge $T_1 \rightarrow R_j$
- assignment edge – directed edge $R_j \rightarrow T_i$


Resource-Allocation Graph Examples

- Model:


- request edge – directed edge $T_1 \rightarrow R_j$
- assignment edge – directed edge $R_j \rightarrow T_i$



Simple Resource
Allocation Graph



Allocation Graph
With Deadlock




Allocation Graph
With Cycle, but
No Deadlock

Basic Facts

- If graph contains no cycles \Rightarrow no deadlock.
- If graph contains a cycle \Rightarrow
 - if only one instance per resource type, then deadlock.
 - if several instances per resource type, possibility of deadlock.

Deadlock Detection Algorithm

- \Rightarrow look for loops
 - More General Deadlock Detection Algorithm
 - Let $[X]$ represent an m-ary vector of non-negative integers (quantities of resources of each type):
 - $[FreeResources]$: Current free resources each type
 - $[Request_x]$: Current requests from thread X
 - $[Alloc_x]$: Current resources held by thread X
 - See if tasks can eventually terminate on their own
- ```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
 done = true
 Foreach node in UNFINISHED {
 if ($[Request_{node}] \leq [Avail]$) {
 remove node from UNFINISHED
 $[Avail] = [Avail] + [Alloc_{node}]$
 done = false
 }
 }
} until(done)
```
- Nodes left in UNFINISHED  $\Rightarrow$  deadlocked



# Four requirements for Deadlock

- Mutual exclusion
  - Only one thread at a time can use a resource.
- Hold and wait
  - Thread holding at least one resource is waiting to acquire additional resources held by other threads
- No preemption
  - Resources are released only voluntarily by the thread holding the resource, after thread is finished with it
- Circular wait
  - There exists a set  $\{T_1, \dots, T_n\}$  of waiting threads
    - $T_1$  is waiting for a resource that is held by  $T_2$
    - $T_2$  is waiting for a resource that is held by  $T_3$
    - ...
    - $T_n$  is waiting for a resource that is held by  $T_1$
- To prevent deadlock, make sure at least one of these conditions does not hold

# How should a system deal with deadlock?

- Three different approaches:
  1. Deadlock avoidance: dynamically delay resource requests so deadlock doesn't happen
  2. Deadlock prevention: write your code in a way that it isn't prone to deadlock
  3. Deadlock recovery: let deadlock happen, and then figure out how to recover from it
- Modern operating systems:
  - Make sure the *system* isn't involved in any deadlock
  - Ignore deadlock in applications
    - “Ostrich Algorithm”

# Deadlock Avoidance

- Idea: When a thread requests a resource, OS checks if it would result in deadlock
  - If not, it grants the resource right away
  - If so, it waits for other threads to release resources

**THIS DOES NOT WORK!!!**

- Example:

**Thread A**

**x.Acquire();**

Blocks... **y.Acquire();**

...

**y.Release();**

**x.Release();**

**Thread B**

**y.Acquire();**

**x.Acquire();**

...

**x.Release();**

**y.Release();**

Wait...


But it's too late...

# Deadlock Avoidance: Three States

- Safe state
  - System can delay resource acquisition to prevent deadlock
- Unsafe state
  - No deadlock yet...  
Deadlock avoidance: prevent system from reaching an *unsafe* state
  - But threads can request resources in a pattern that unavoidably leads to deadlock
- Deadlocked state
  - There exists a deadlock in the system
  - Also considered “unsafe”



# Safe, Unsafe , Deadlock State



# Deadlock Avoidance

- Idea: When a thread requests a resource, OS checks if it would result in ~~deadlock~~ **an unsafe state**
  - If not, it grants the resource right away
  - If so, it waits for other threads to release resources
- Example:

## Thread A

```
x.Acquire();
y.Acquire();
...
y.Release();
x.Release();
```

## Thread B

```
y.Acquire();
x.Acquire();
...
x.Release();
y.Release();
```

Wait until  
Thread A  
releases the  
mutex

# Deadlock Avoidance with Resource Reservation

- Threads provide advance information about the maximum resources they may need during execution
- Define a sequence of threads  $\{T_1, T_2, \dots T_n\}$  as *safe* if for each  $T_i$  the resources that  $T_i$  can still request can be satisfied by the currently available resources plus the resources held by all  $T_j, j < i$ .
- A *safe state* is a state in which there is a safe sequence for the threads.
- An unsafe state is not equivalent to deadlock, it just may lead to deadlock, since some threads might not actually use the maximum resources they have declared.
- Grant a resource to a thread if the new state is safe
- If the new state is unsafe, the thread must wait even if the resource is currently available.
- This algorithm ensures no circular-wait condition exists.


# Example

Threads **t1**, **t2**, and **t3** are competing for 12 tape drives.

- Currently, 11 drives are allocated to the threads, leaving 1 available.
- The current state is *safe* (there exists a safe sequence, {t1, t2, t3} where all threads may obtain their maximum number of resources without waiting)
  - t1 can complete with the current resource allocation
  - t2 can complete with its current resources, plus all of t1's resources, and the unallocated tape drive.
- t3 can complete with all its current resources, all of t1 and t2's resources, and the unallocated tape drive.

|                | max<br>need | in use | could<br>want |
|----------------|-------------|--------|---------------|
| t <sub>1</sub> | 4           | 3      | 1             |
| t <sub>2</sub> | 8           | 4      | 4             |
| t <sub>3</sub> | 12          | 4      | 8             |


# Banker's Algorithm



- idea:
  - State maximum resource needs in advance
  - Allow particular thread to proceed if:  
 $(\text{available resources} - \text{\#requested}) \geq \text{max remaining that might be needed by any thread}$
- Banker's algorithm
  - Allocate resources dynamically
    - Evaluate each request and grant if some ordering of threads is still deadlock free afterward
    - Keeps system in a "SAFE" state, i.e. there exists a sequence  $\{T_1, T_2, \dots, T_n\}$  with  $T_1$  requesting all remaining resources, finishing, then  $T_2$  requesting all remaining resources, etc..
  - Algorithm allows the sum of maximum resource needs of all current threads to be greater than total resources



# Banker's Algorithm



- Technique: pretend each request is granted, then run deadlock detection algorithm, substitute

$$([Request_{node}] \leq [Avail]) \rightarrow ([Max_{node}] - [Alloc_{node}] \leq [Avail])$$

|                        |                                     |
|------------------------|-------------------------------------|
| [FreeResources]:       | Current free resources each type    |
| [Alloc <sub>x</sub> ]: | Current resources held by thread X  |
| [Max <sub>x</sub> ]:   | Max resources requested by thread X |


```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
 done = true
 Foreach node in UNFINISHED {
 if ([Maxnode] - [Allocnode] <= [Avail]) {
 remove node from UNFINISHED
 [Avail] = [Avail] + [Allocnode]
 done = false
 }
 }
} until(done)
```

# Example of Banker's Algorithm

- 5 Process/Treads  $T_0$  through  $T_4$ ;  
3 resource types:  
     $A$  (10 instances),  $B$  (5instances), and  $C$  (7 instances).

The content of the matrix *Need* is defined to be  
*Max – Allocation*.

## Example of Banker's algorithm



Non-allocated  
resource instances.

| Available |   |   |
|-----------|---|---|
| A         | B | C |
| 3         | 3 | 2 |

Allocated resource  
instances.

|       | Allocation |   |   |
|-------|------------|---|---|
|       | A          | B | C |
| $T_0$ | 0          | 1 | 0 |
| $T_1$ | 2          | 0 | 0 |
| $T_2$ | 3          | 0 | 2 |
| $T_3$ | 2          | 1 | 1 |
| $T_4$ | 0          | 0 | 2 |



# Example of Banker's algorithm

|                | Max |   |   |
|----------------|-----|---|---|
|                | A   | B | C |
| T <sub>0</sub> | 7   | 5 | 3 |
| T <sub>1</sub> | 3   | 2 | 2 |
| T <sub>2</sub> | 9   | 0 | 2 |
| T <sub>3</sub> | 2   | 2 | 2 |
| T <sub>4</sub> | 4   | 3 | 3 |

|                | Need |   |   |
|----------------|------|---|---|
|                | A    | B | C |
| T <sub>0</sub> | 7    | 4 | 3 |
| T <sub>1</sub> | 1    | 2 | 2 |
| T <sub>2</sub> | 6    | 0 | 0 |
| T <sub>3</sub> | 0    | 1 | 1 |
| T <sub>4</sub> | 4    | 3 | 1 |

|                | Allocation |   |   |
|----------------|------------|---|---|
|                | A          | B | C |
| T <sub>0</sub> | 0          | 1 | 0 |
| T <sub>1</sub> | 2          | 0 | 0 |
| T <sub>2</sub> | 3          | 0 | 2 |
| T <sub>3</sub> | 2          | 1 | 1 |
| T <sub>4</sub> | 0          | 0 | 2 |

| Available |   |   |
|-----------|---|---|
| A         | B | C |
| 3         | 3 | 2 |

$$\text{Need} = \text{Max} - \text{Allocation}$$

# Example of Banker's algorithm

|                | Max |   |   |
|----------------|-----|---|---|
|                | A   | B | C |
| T <sub>0</sub> | 7   | 5 | 3 |
| T <sub>1</sub> | 3   | 2 | 2 |
| T <sub>2</sub> | 9   | 0 | 2 |
| T <sub>3</sub> | 2   | 2 | 2 |
| T <sub>4</sub> | 4   | 3 | 3 |

|                | Need |   |   |
|----------------|------|---|---|
|                | A    | B | C |
| T <sub>0</sub> | 7    | 4 | 3 |
| T <sub>1</sub> | 1    | 2 | 2 |
| T <sub>2</sub> | 6    | 0 | 0 |
| T <sub>3</sub> | 0    | 1 | 1 |
| T <sub>4</sub> | 4    | 3 | 1 |

| Available |   |   | Available |   |   |
|-----------|---|---|-----------|---|---|
| A         | B | C | A         | B | C |
| 3         | 3 | 2 | 2         | 1 | 0 |

|                | Allocation |   |   | Allocation |   |   |
|----------------|------------|---|---|------------|---|---|
|                | A          | B | C | A          | B | C |
| T <sub>0</sub> | 0          | 1 | 0 | 0          | 1 | 0 |
| T <sub>1</sub> | 2          | 0 | 0 | 3          | 2 | 2 |
| T <sub>2</sub> | 3          | 0 | 2 | 3          | 0 | 2 |
| T <sub>3</sub> | 2          | 1 | 1 | 2          | 1 | 1 |
| T <sub>4</sub> | 0          | 0 | 2 | 0          | 0 | 2 |

Allocation for T<sub>1</sub> increases and now T<sub>1</sub> holds its max (3, 2, 2).

# Example of Banker's algorithm

|                | Max |   |   |
|----------------|-----|---|---|
|                | A   | B | C |
| T <sub>0</sub> | 7   | 5 | 3 |
| T <sub>1</sub> | 3   | 2 | 2 |
| T <sub>2</sub> | 9   | 0 | 2 |
| T <sub>3</sub> | 2   | 2 | 2 |
| T <sub>4</sub> | 4   | 3 | 3 |

|                | Need |   |   |
|----------------|------|---|---|
|                | A    | B | C |
| T <sub>0</sub> | 7    | 4 | 3 |
| T <sub>1</sub> | 1    | 2 | 2 |
| T <sub>2</sub> | 6    | 0 | 0 |
| T <sub>3</sub> | 0    | 1 | 1 |
| T <sub>4</sub> | 4    | 3 | 1 |

|                | Allocation |   |   |
|----------------|------------|---|---|
|                | A          | B | C |
| T <sub>0</sub> | 0          | 1 | 0 |
| T <sub>1</sub> | 2          | 0 | 0 |
| T <sub>2</sub> | 3          | 0 | 2 |
| T <sub>3</sub> | 2          | 1 | 1 |
| T <sub>4</sub> | 0          | 0 | 2 |

| Available |   |   |
|-----------|---|---|
| A         | B | C |
| 3         | 3 | 2 |

$$\text{Need} = \text{Max} - \text{Allocation}$$

# Example of Banker's algorithm

$Available_{before} + Allocation[T_1]_{before}$

$$(3, 3, 2) + (2, 0, 0) = (5, 3, 2)$$

**Before**

**Under**

**After**

| Available |   |   | Available |   |   | Available |   |   |
|-----------|---|---|-----------|---|---|-----------|---|---|
| A         | B | C | A         | B | C | A         | B | C |
| 3         | 3 | 2 | 2         | 1 | 0 | 5         | 3 | 2 |

|                | Max |   |   |
|----------------|-----|---|---|
|                | A   | B | C |
| T <sub>0</sub> | 7   | 5 | 3 |
| T <sub>1</sub> | 3   | 2 | 2 |
| T <sub>2</sub> | 9   | 0 | 2 |
| T <sub>3</sub> | 2   | 2 | 2 |
| T <sub>4</sub> | 4   | 3 | 3 |

|                | Need |   |   |
|----------------|------|---|---|
|                | A    | B | C |
| T <sub>0</sub> | 7    | 4 | 3 |
| T <sub>1</sub> | 1    | 2 | 2 |
| T <sub>2</sub> | 6    | 0 | 0 |
| T <sub>3</sub> | 0    | 1 | 1 |
| T <sub>4</sub> | 4    | 3 | 1 |

|                | Allocation |   |   | Allocation |   |   | Allocation |   |   |
|----------------|------------|---|---|------------|---|---|------------|---|---|
|                | A          | B | C | A          | B | C | A          | B | C |
| T <sub>0</sub> | 0          | 1 | 0 | 0          | 1 | 0 | 0          | 1 | 0 |
| T <sub>1</sub> | 2          | 0 | 0 | 3          | 2 | 2 | 0          | 0 | 0 |
| T <sub>2</sub> | 3          | 0 | 2 | 3          | 0 | 2 | 3          | 0 | 2 |
| T <sub>3</sub> | 2          | 1 | 1 | 2          | 1 | 1 | 2          | 1 | 1 |
| T <sub>4</sub> | 0          | 0 | 2 | 0          | 0 | 2 | 0          | 0 | 2 |

## Example of Banker's algorithm

|                | Need |   |   | Allocation |   |   |      |
|----------------|------|---|---|------------|---|---|------|
|                | A    | B | C | A          | B | C |      |
| T <sub>0</sub> | 7    | 4 | 3 | 0          | 1 | 0 | Done |
| T <sub>1</sub> | 1    | 2 | 2 | 2          | 0 | 0 | Done |
| T <sub>2</sub> | 6    | 0 | 0 | 3          | 0 | 2 | Done |
| T <sub>3</sub> | 0    | 1 | 1 | 2          | 1 | 1 | Done |
| T <sub>4</sub> | 4    | 3 | 1 | 0          | 0 | 2 | Done |

| Step | Available |   |   | Done           | Choice         |
|------|-----------|---|---|----------------|----------------|
|      | A         | B | C |                |                |
| 1    | 3         | 3 | 2 | -              | T <sub>1</sub> |
| 2    | 5         | 3 | 2 | T <sub>1</sub> | T <sub>3</sub> |
| 3    | 7         | 4 | 3 | T <sub>3</sub> | T <sub>4</sub> |
| 4    | 7         | 4 | 5 | T <sub>4</sub> | T <sub>2</sub> |
| 5    | 10        | 4 | 7 | T <sub>2</sub> | T <sub>0</sub> |
| 6    | 10        | 5 | 7 | T <sub>0</sub> | ★              |

sequence  $\langle T_1, T_3, T_4, T_2, T_0 \rangle$  that allowed all tasks to get their resources - **state is safe!**

## Example of Banker's algorithm

|                | Need |   |   | Allocation |   |   |      |
|----------------|------|---|---|------------|---|---|------|
|                | A    | B | C | A          | B | C |      |
| T <sub>0</sub> | 7    | 2 | 3 | 0          | 3 | 0 | Done |
| T <sub>1</sub> | 1    | 2 | 2 | 2          | 0 | 0 | Done |
| T <sub>2</sub> | 6    | 0 | 0 | 3          | 0 | 2 | Done |
| T <sub>3</sub> | 0    | 1 | 1 | 2          | 1 | 1 | Done |
| T <sub>4</sub> | 4    | 3 | 1 | 0          | 0 | 2 | Done |

| Step | Available |   |   | Done           | Choice         |
|------|-----------|---|---|----------------|----------------|
|      | A         | B | C |                |                |
| 1    | 3         | 1 | 2 | -              | T <sub>3</sub> |
| 2    | 5         | 2 | 3 | T <sub>3</sub> | T <sub>1</sub> |
| 3    | 7         | 2 | 3 | T <sub>1</sub> | T <sub>0</sub> |
| 4    | 7         | 5 | 3 | T <sub>0</sub> | T <sub>2</sub> |
| 5    | 10        | 5 | 5 | T <sub>2</sub> | T <sub>4</sub> |
| 6    | 10        | 5 | 7 | T <sub>4</sub> | ★              |

sequence < T<sub>3</sub>, T<sub>1</sub>, T<sub>0</sub>, T<sub>2</sub>, T<sub>4</sub> > that allowed all tasks to get their resources - **state is safe!**

# Deadlock Detection

- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

# Deadlock Detection

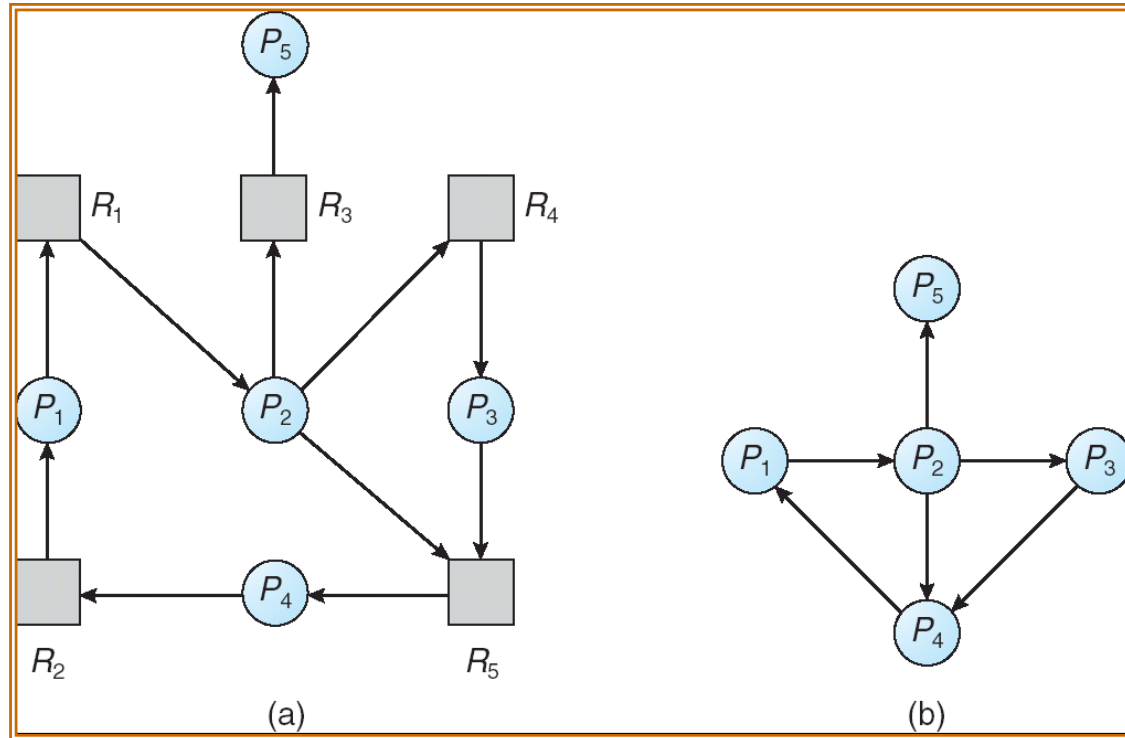
- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme



# Single Instance of Each Resource Type

- Maintain *wait-for* graph
  - Nodes are processes.
  - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$ .
- Periodically invoke an algorithm that searches for a cycle in the graph. If there is a cycle, there exists a deadlock.
- An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph.

# Resource-Allocation Graph and Wait-for Graph



Resource-Allocation Graph

Corresponding wait-for graph

# Several Instances of a Resource Type

- **Available:** A vector of length  $m$  indicates the number of available resources of each type.
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process.
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If  $Request[i_j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

# Example of Detection Algorithm

- Five processes  $P_0$  through  $P_4$ ; three resource types A (7 instances), B (2 instances), and C (6 instances).
- Snapshot at time  $T_0$ :

|       | <u>Allocation</u> |   |   | <u>Request</u> |   |   | <u>Available</u> |   |   |
|-------|-------------------|---|---|----------------|---|---|------------------|---|---|
|       | A                 | B | C | A              | B | C | A                | B | C |
| $P_0$ | 0                 | 1 | 0 | 0              | 0 | 0 | 0                | 0 | 0 |
| $P_1$ | 2                 | 0 | 0 | 2              | 0 | 2 |                  |   |   |
| $P_2$ | 3                 | 0 | 3 | 0              | 0 | 0 |                  |   |   |
| $P_3$ | 2                 | 1 | 1 | 1              | 0 | 0 |                  |   |   |
| $P_4$ | 0                 | 0 | 2 | 0              | 0 | 2 |                  |   |   |

- Sequence  $\langle P_0, P_2, P_3, P_1, P_4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$ .

## Example (Cont.)

- $P_2$  requests an additional instance of type C.

|       | <u>Request</u> |   |   |
|-------|----------------|---|---|
|       | A              | B | C |
| $P_0$ | 0              | 0 | 0 |
| $P_1$ | 2              | 0 | 1 |
| $P_2$ | 0              | 0 | 1 |
| $P_3$ | 1              | 0 | 0 |
| $P_4$ | 0              | 0 | 2 |

- State of system?
  - Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes; requests.
  - Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$ .

# Detection Algorithm

1. Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively Initialize:
  - (a) *Work* = *Available*
  - (b) For  $i = 1, 2, \dots, n$ , if  $Allocation_i \neq 0$ , then  $Finish[i] = false$ ; otherwise,  $Finish[i] = true$ .
2. Find an index  $i$  such that both:
  - (a)  $Finish[i] == false$
  - (b)  $Request_i \leq Work$

If no such  $i$  exists, go to step 4.

# Detection Algorithm (Cont.)

3.  $Work = Work + Allocation_i$   
 $Finish[i] = true$   
go to step 2.
4. If  $Finish[i] == false$ , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if  $Finish[i] == false$ , then  $P_i$  is deadlocked.

Algorithm requires an order of  $O(m \times n^2)$  operations to detect whether the system is in deadlocked state.